

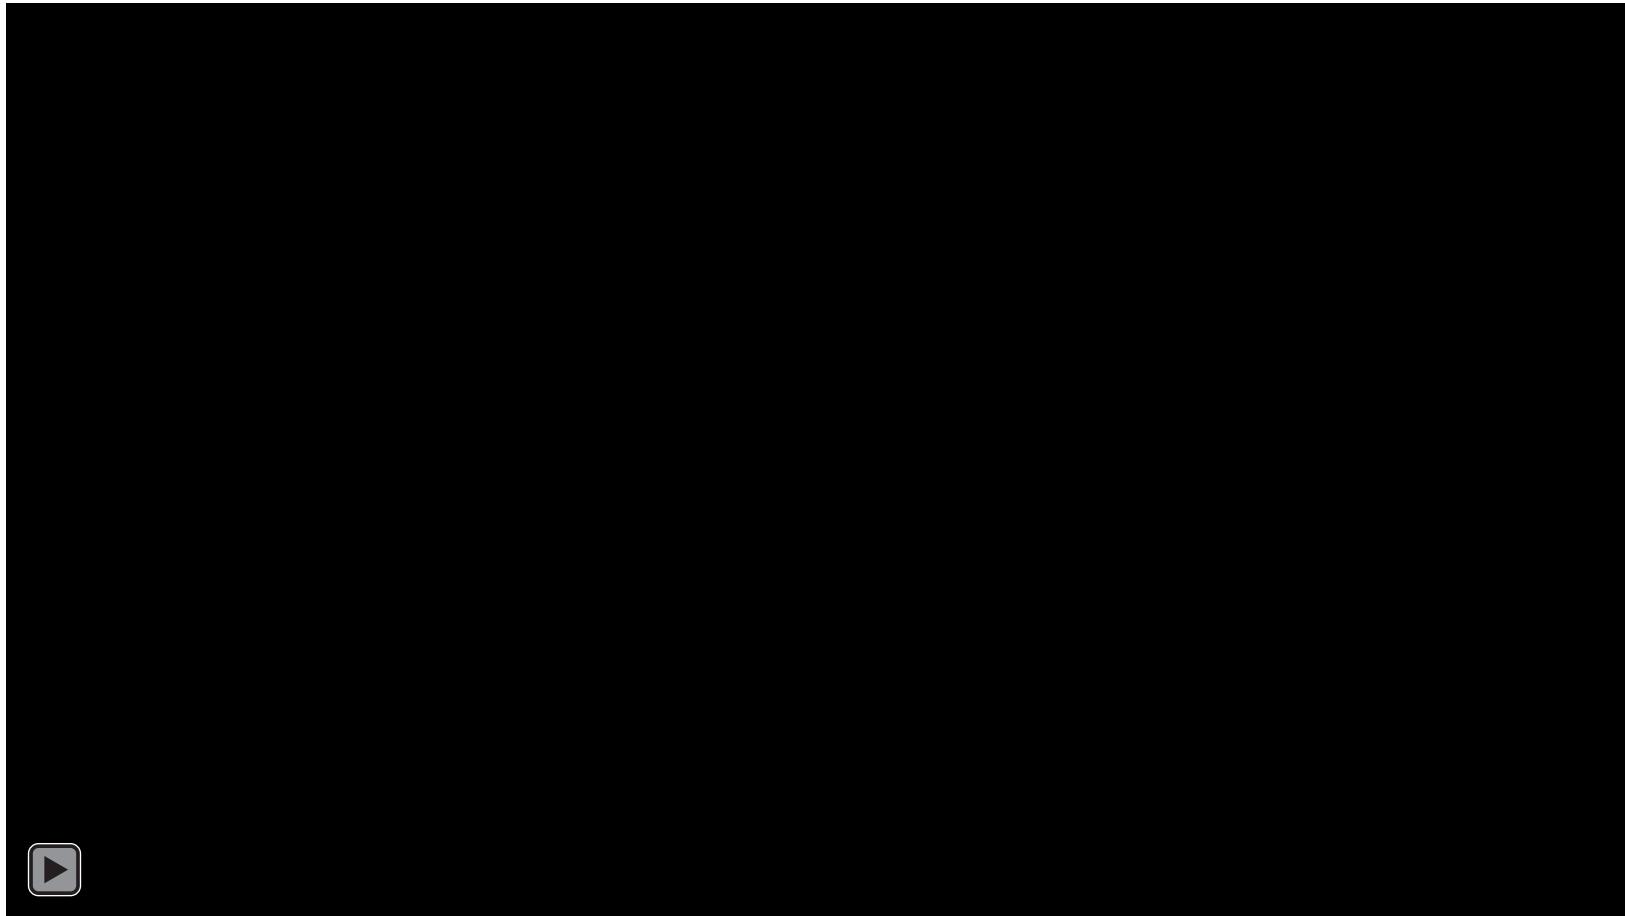
"HIS PATH-PLANNING MAY BE  
SUB-OPTIMAL, BUT IT'S GOT FLAIR."

# Robot Intelligence: Planning

# Assignments

- Due Today, 8/19
  - Read Ch. 3 in Russel & Norvig
- Due Monday, 8/24
  - Read Ch. 4 in Russel & Norvig
- Due Wednesday, 8/26
  - Pset1 (released on Canvas by 3:15 pm on Wednesday, 8/19)

# Research Opportunity



# Outline:

## Problem Solving as State Space Search

- 
- Problem Formulation (Modeling)
    - Problem solving as state space search
  - Formal Representation
    - Graphs and search trees
  - Reasoning Algorithms
    - Depth and breadth-first search

# Search is the basic building block

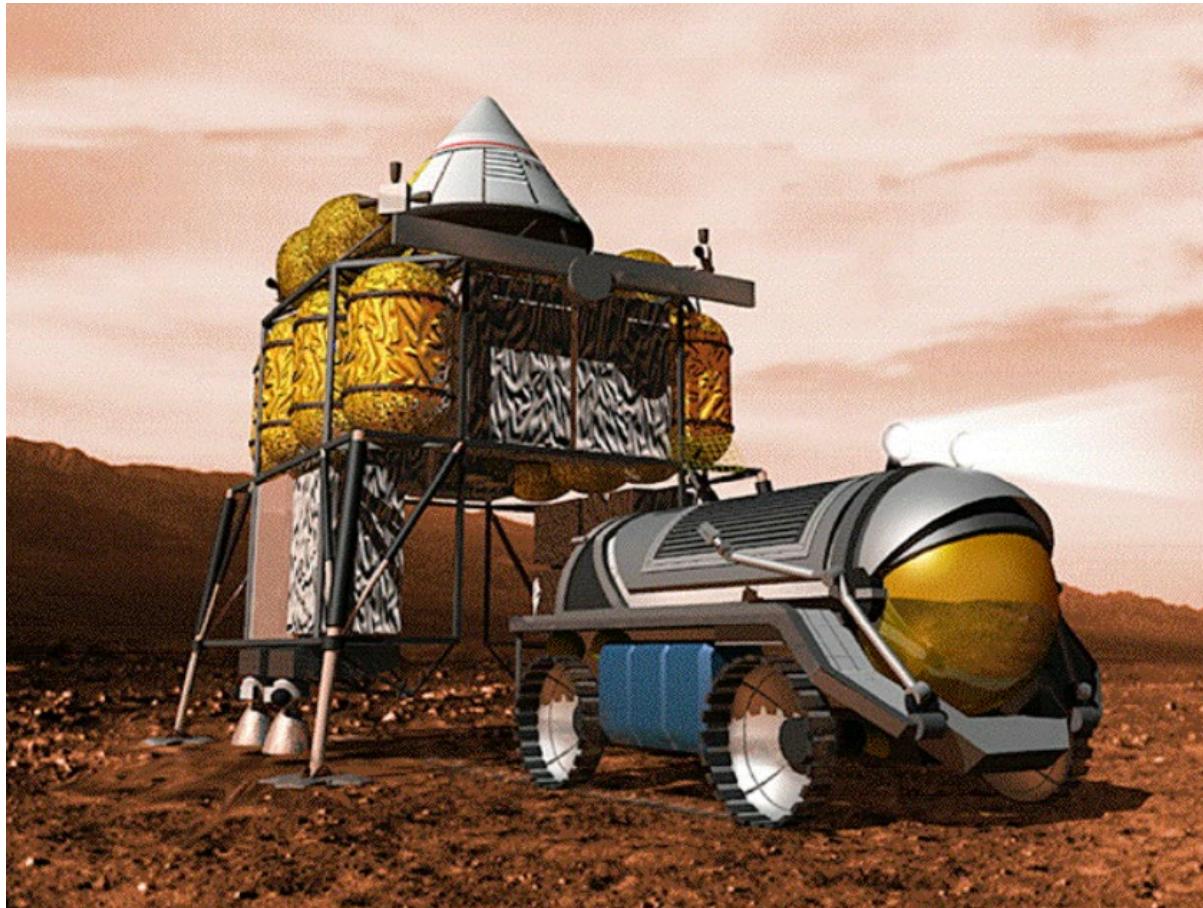
## Operations/Task Planning:

- Activity Planning
- Diagnosis
- Repair
- Scheduling
- Resource Allocation

## Mobility/Motion Planning:

- Path Planning
- Localization
- Map Building
- Control Trajectory Design

# Example: Outpost Logistics Planning



*Image produced for NASA by John Frassanito and Associates*

# Example: Outpost Logistics Planning

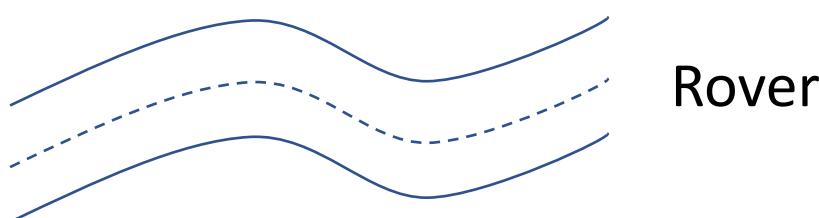
- Can the astronaut get its supplies safely across a Martian crevasse?

Astronaut

Goose

Grain

Fox



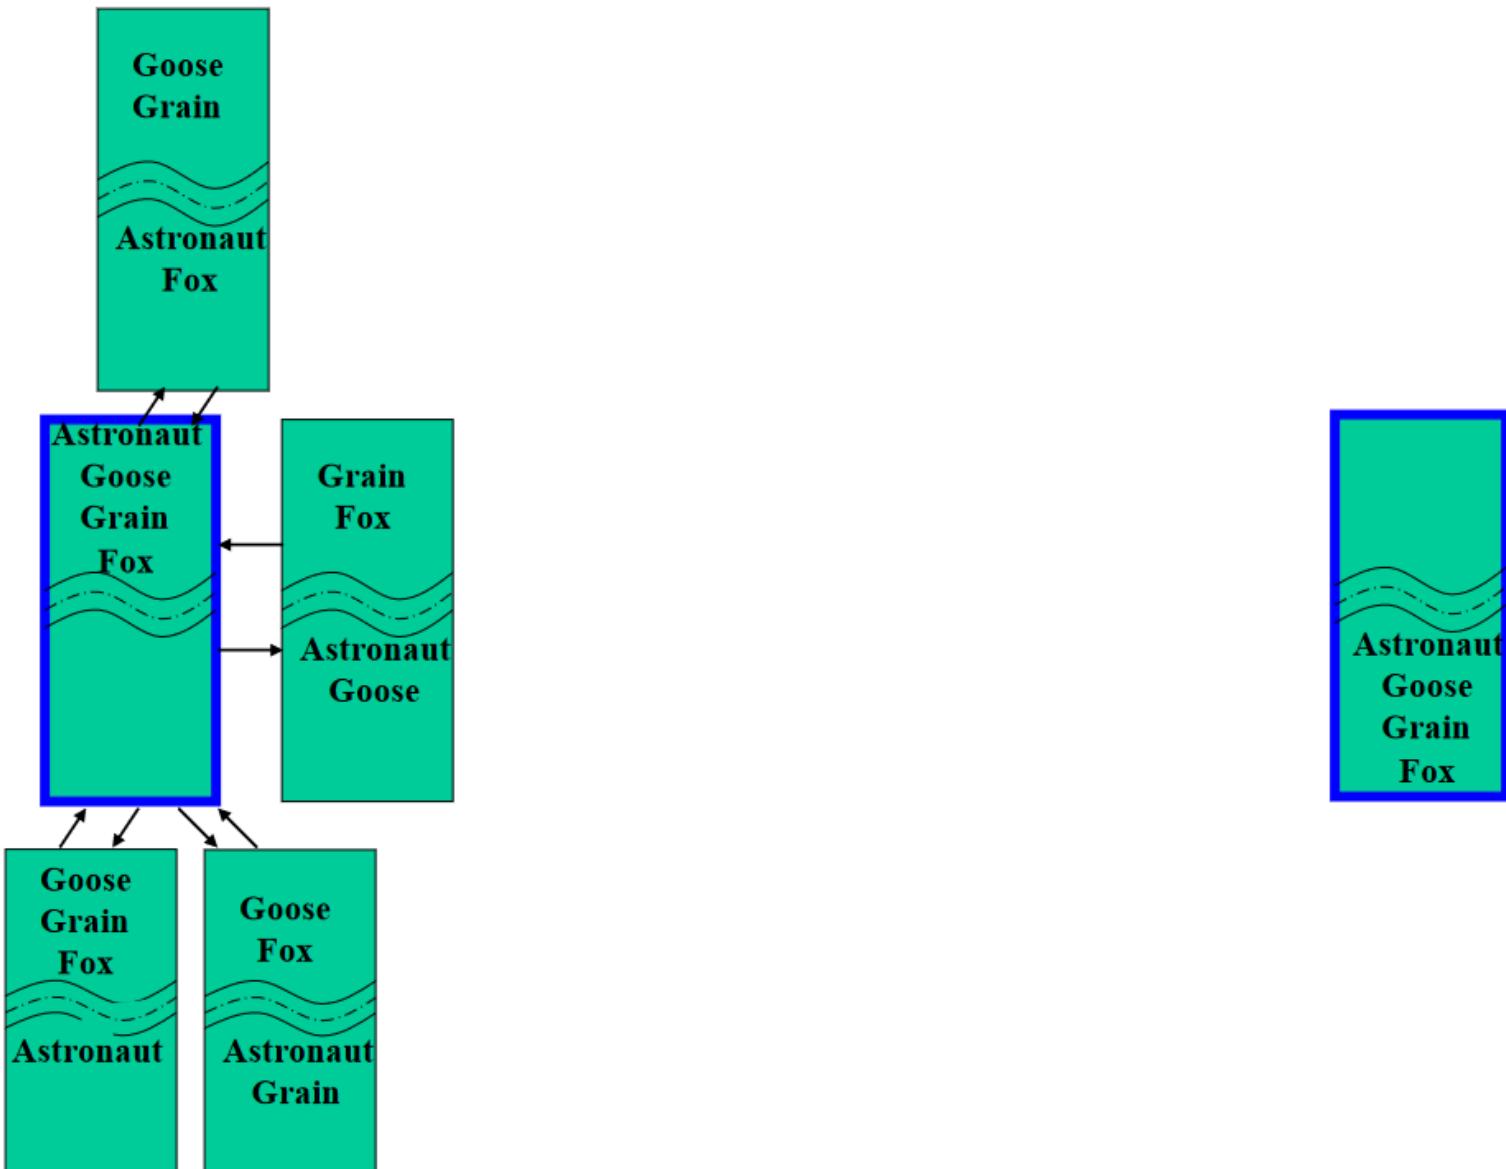
- Astronaut + 1 item allowed in the rover
- Goose alone eats Grain
- Fox alone eats Goose

- Early AI: What are the universal problem solving methods?

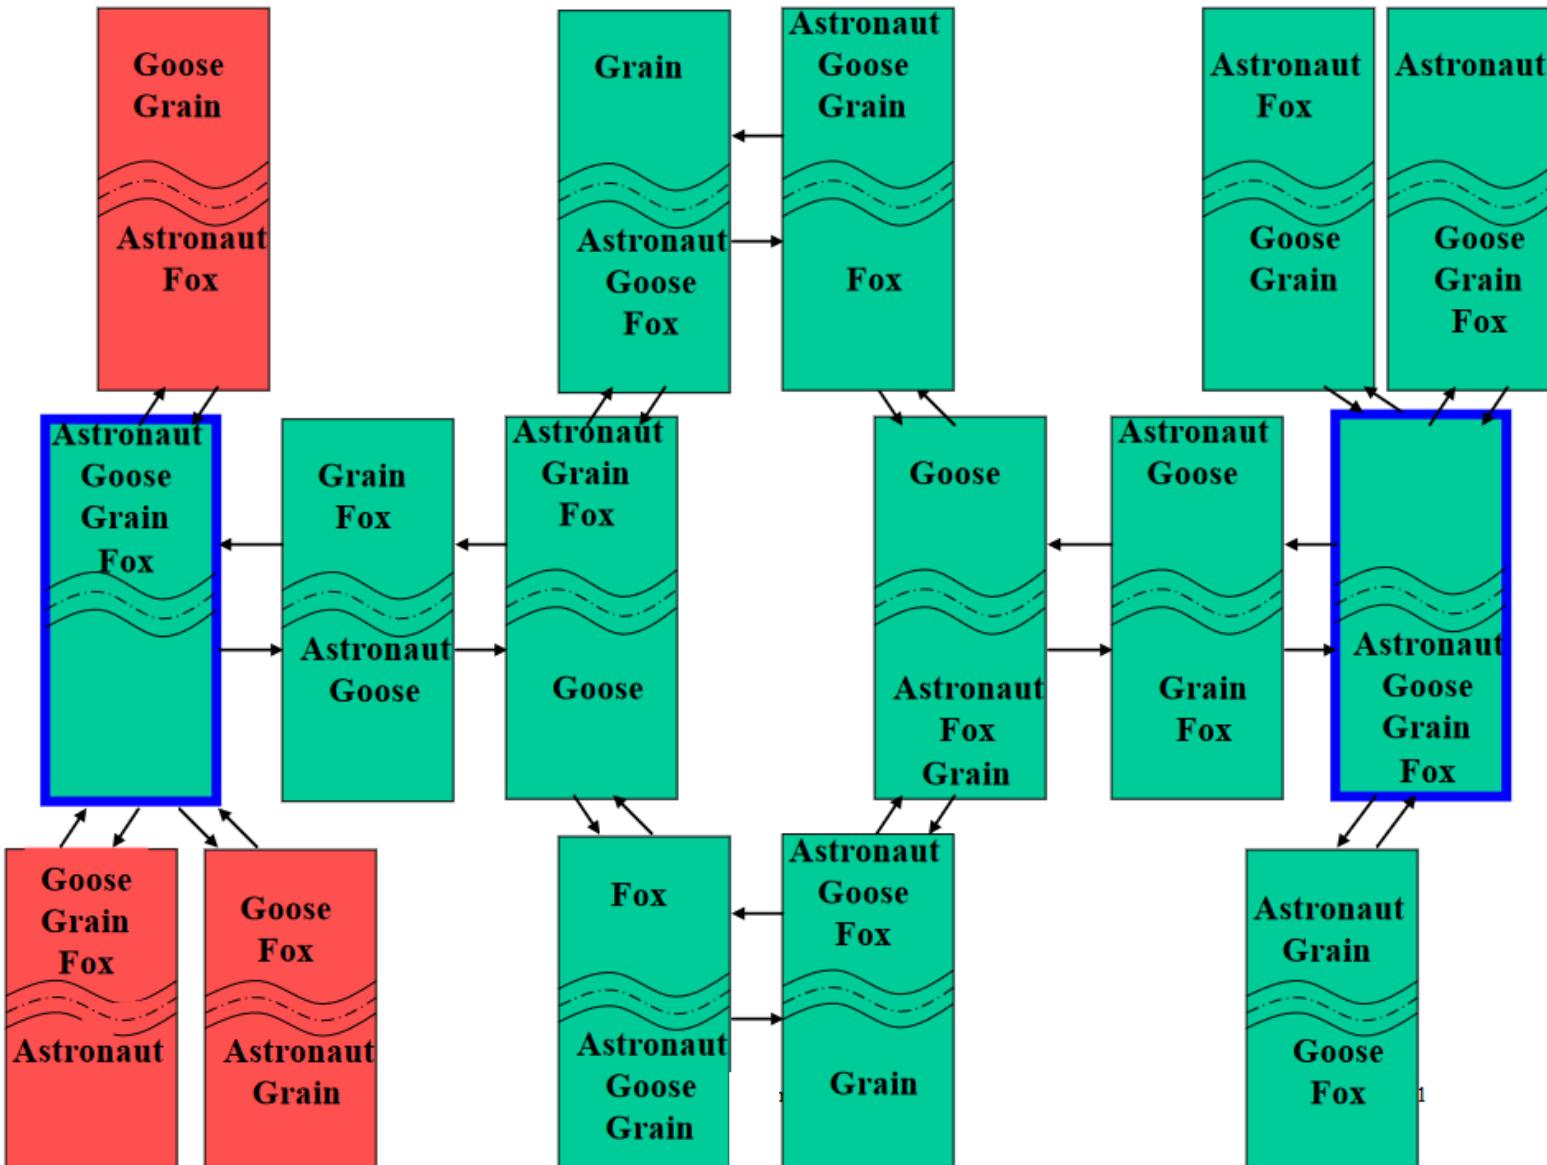
Simple Trivial

# Problem Solving as State Space Search

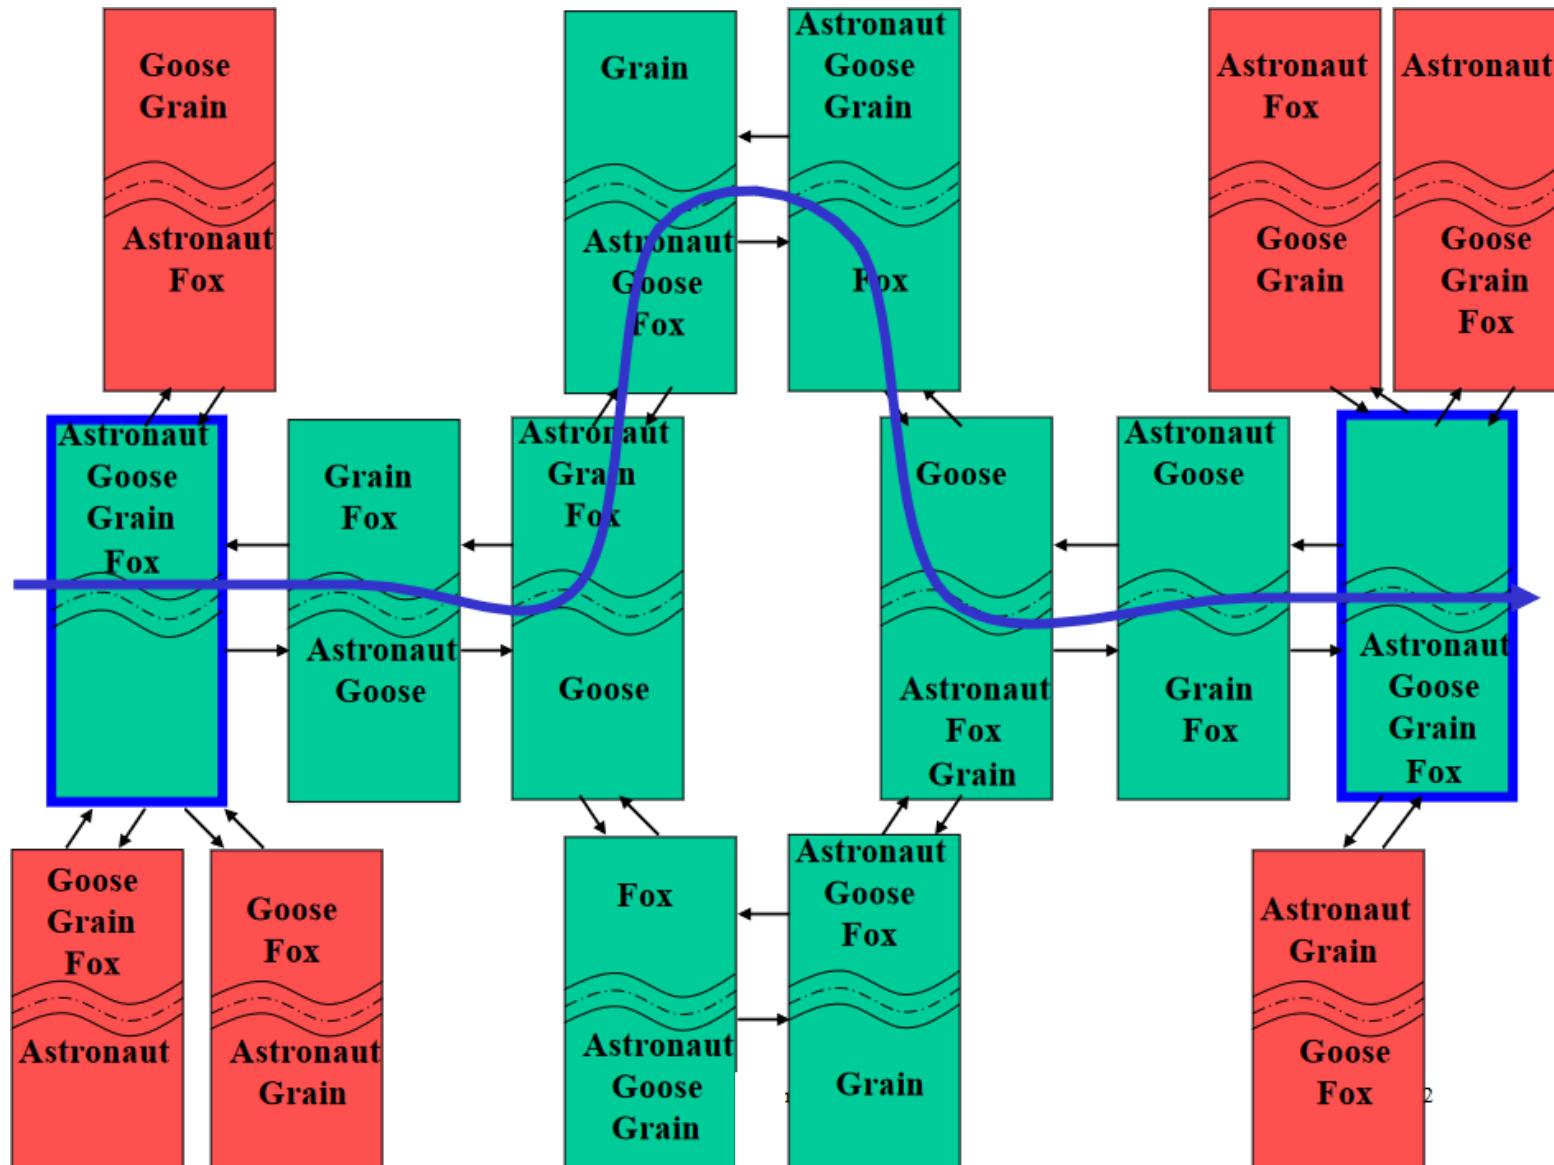
- Formulate Goal
  - State
    - Astronaut, Fox, Goose, & Grain below crevasse
- Formulate Problem
  - States
    - Astronaut, Fox, Goose, & Grain above or below the crevasse
  - Operators
    - Move: Astronaut drives rover and 1 or 0 items to the other side of the crevasse
  - Initial State
    - Astronaut, Fox, Goose, & Grain above crevasse
- Generate Solution
  - Sequence of Operators (or States)
    - Move(goose,astronaut), Move(astronaut), ...



Credit: Brian Williams



Credit: Brian Williams



Credit: Brian Williams

# Formulation Example: 8-Puzzle

5	4	
6	1	8
7	3	2

Start

1	2	3
8		4
7	6	5

Goal

- States: integer location for each tile AND ...???
- Operators: move empty square up, down, left, right
- Initial and Goal States: as shown above

# Languages for Express States and Operators for Complex Tasks



Image Source: NASA

# Example: STRIPS Planning Language

Stanford Research Institute Problem Solver

## Initial State:

```
(and (hose a)
     (clamp b)
     (hydroxide-unit c)
     (on-table a)
     (on-table b)
     (on-table c)
     (clear a)
     (clear b)
     (clear c)
     (empty arm))
```

## Goal (Partial) State:

```
(and (connected a b)
     (connected b c))
```

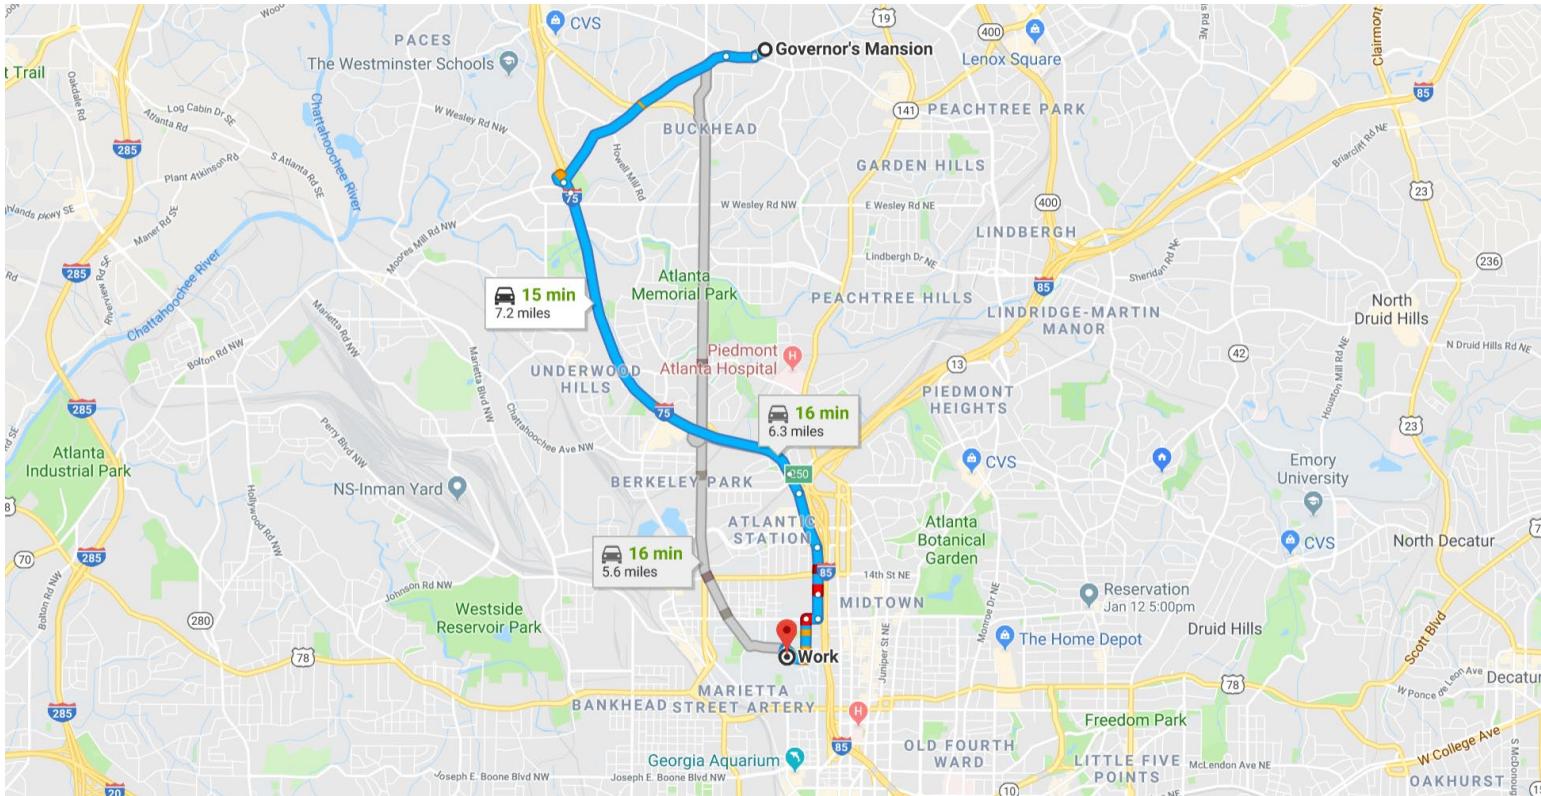
## Operators:

Precondition: (and (clear hose)  
(on-table hose)  
(empty arm))

Pickup hose

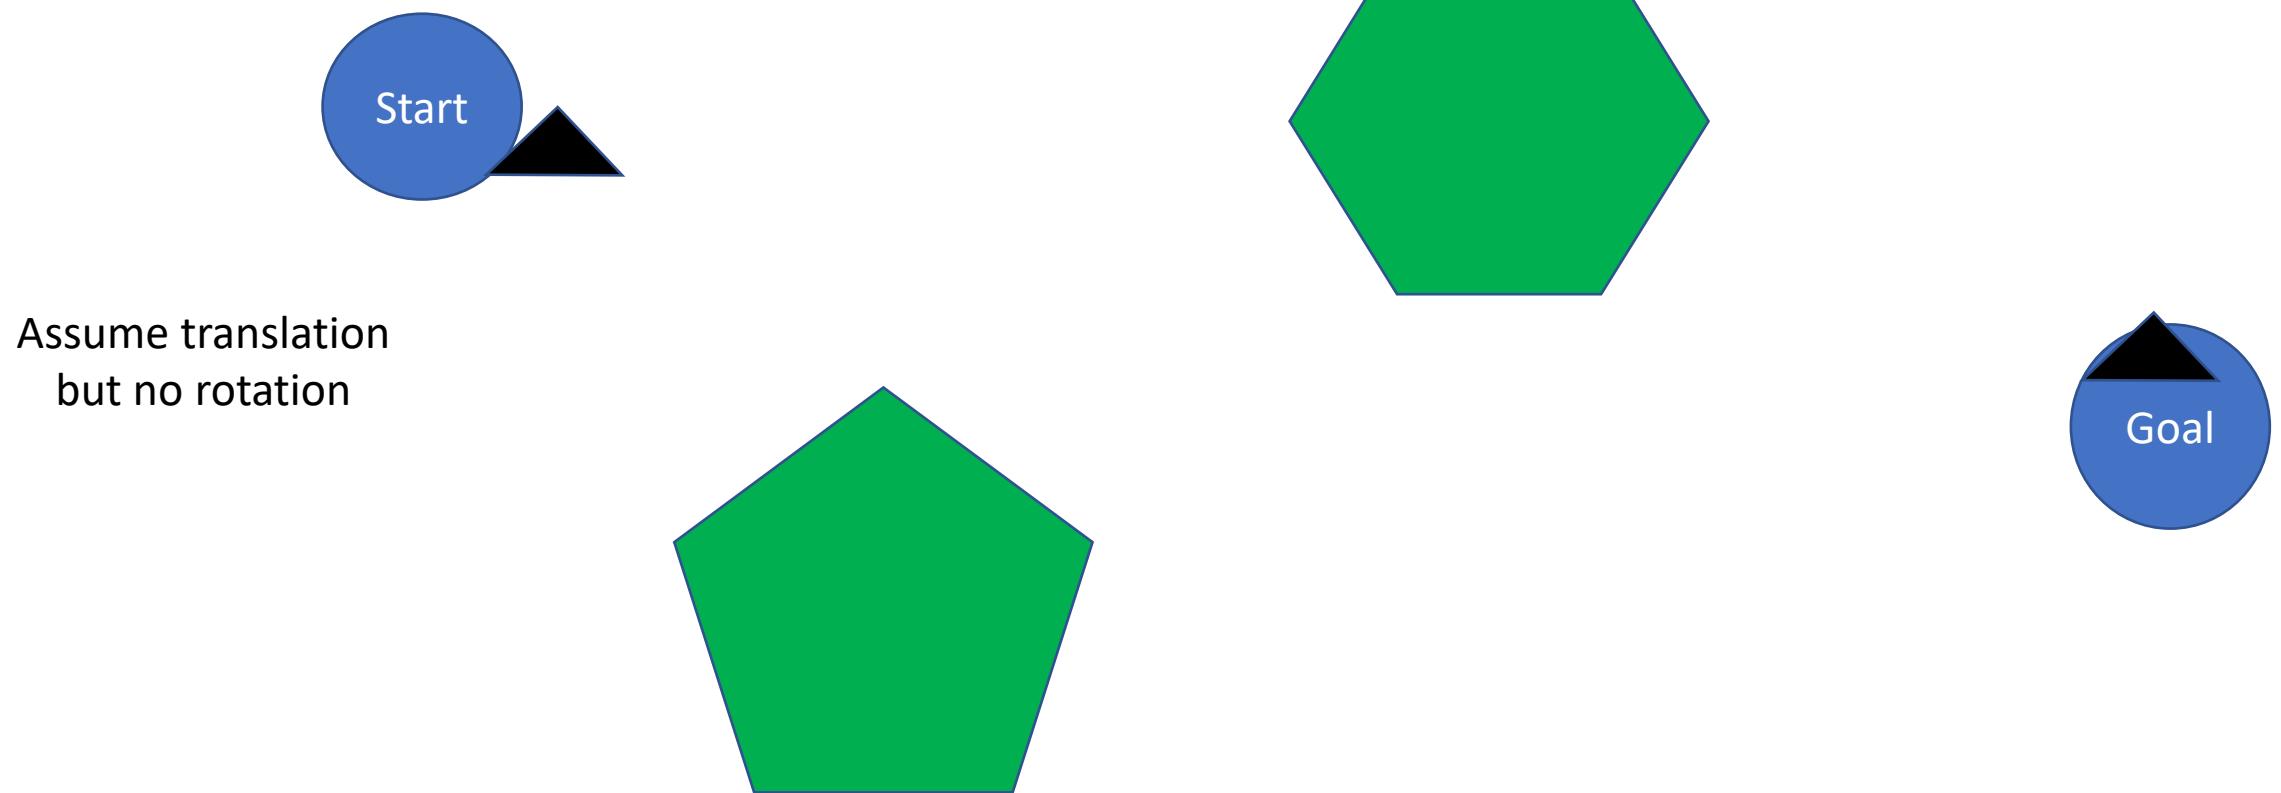
Effect: (and (not (clear hose))  
(not (on-table hose))  
(not (empty arm))  
(holding arm hose))

# Problem: Find a route from home to GaTech

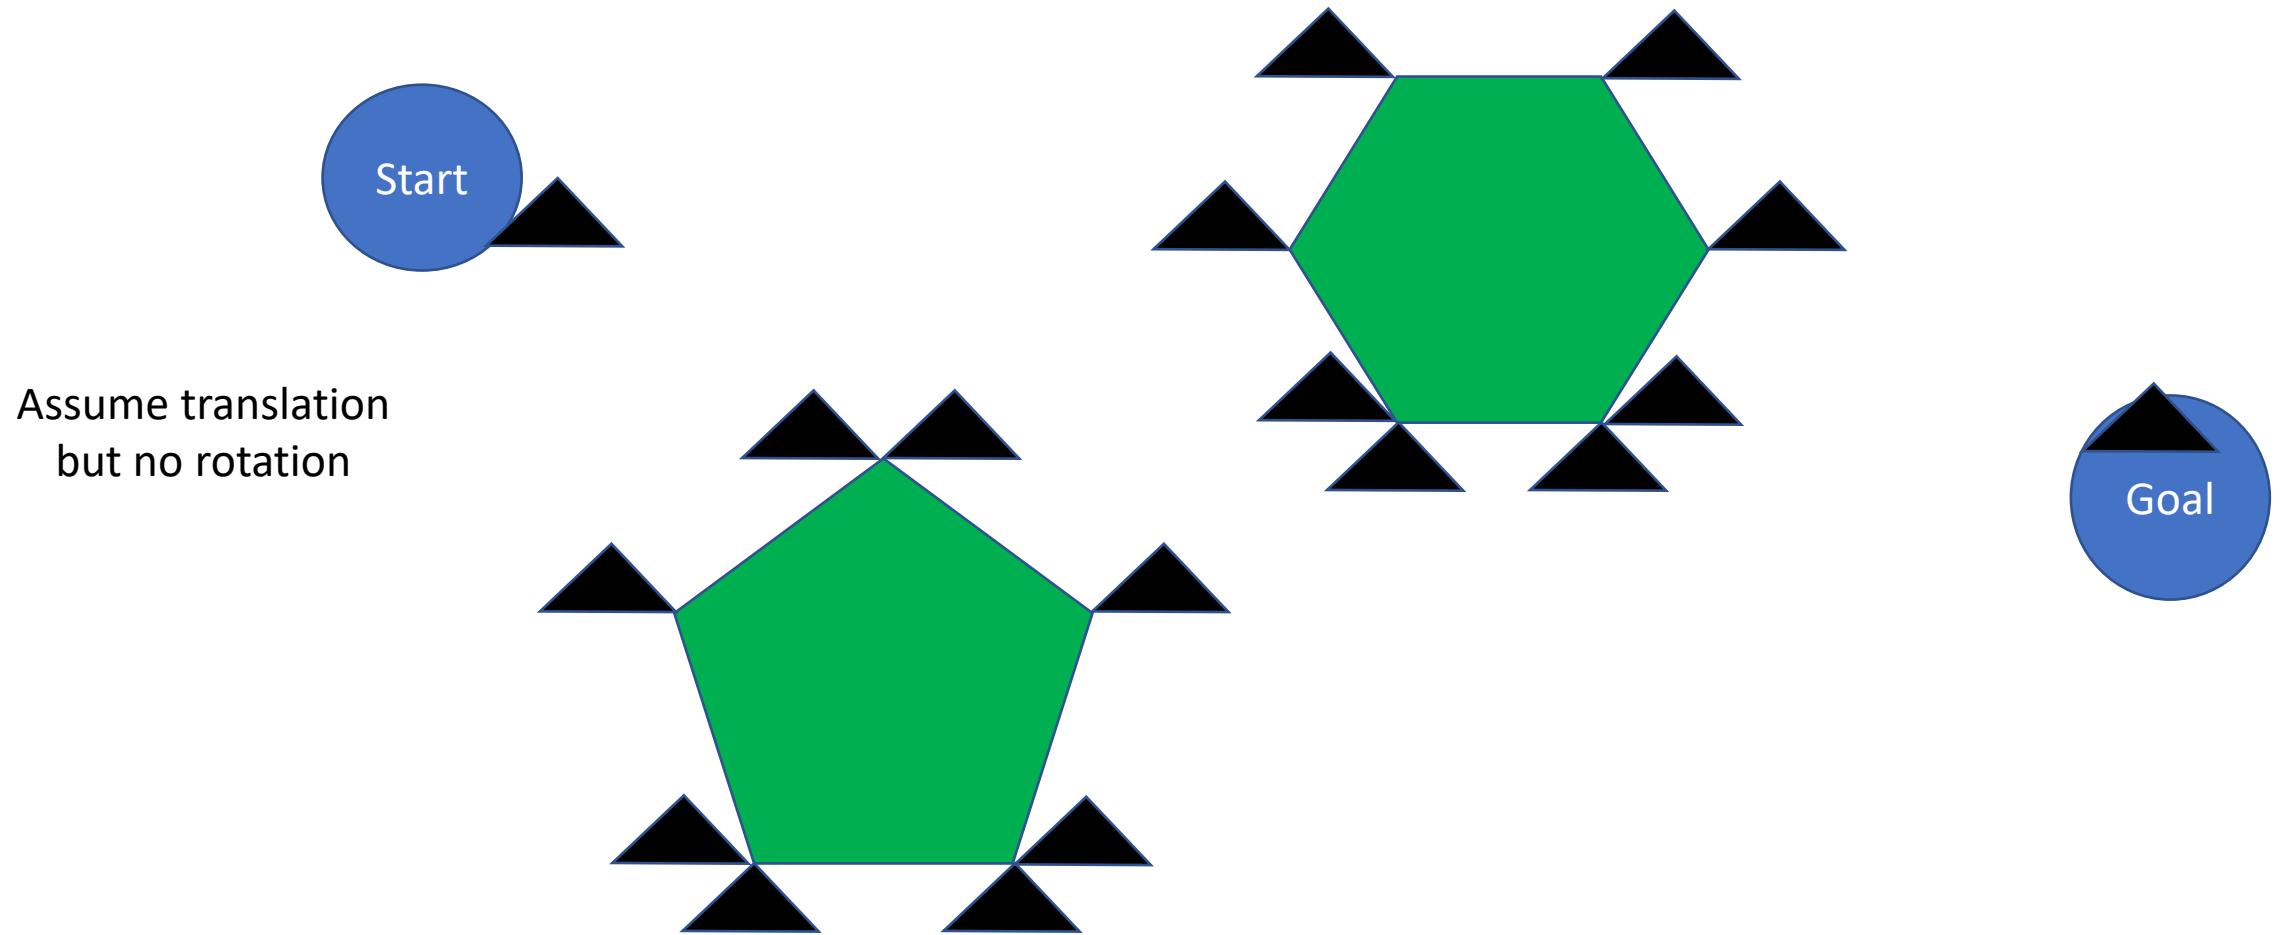


States? Operators? Initial and Goal State?

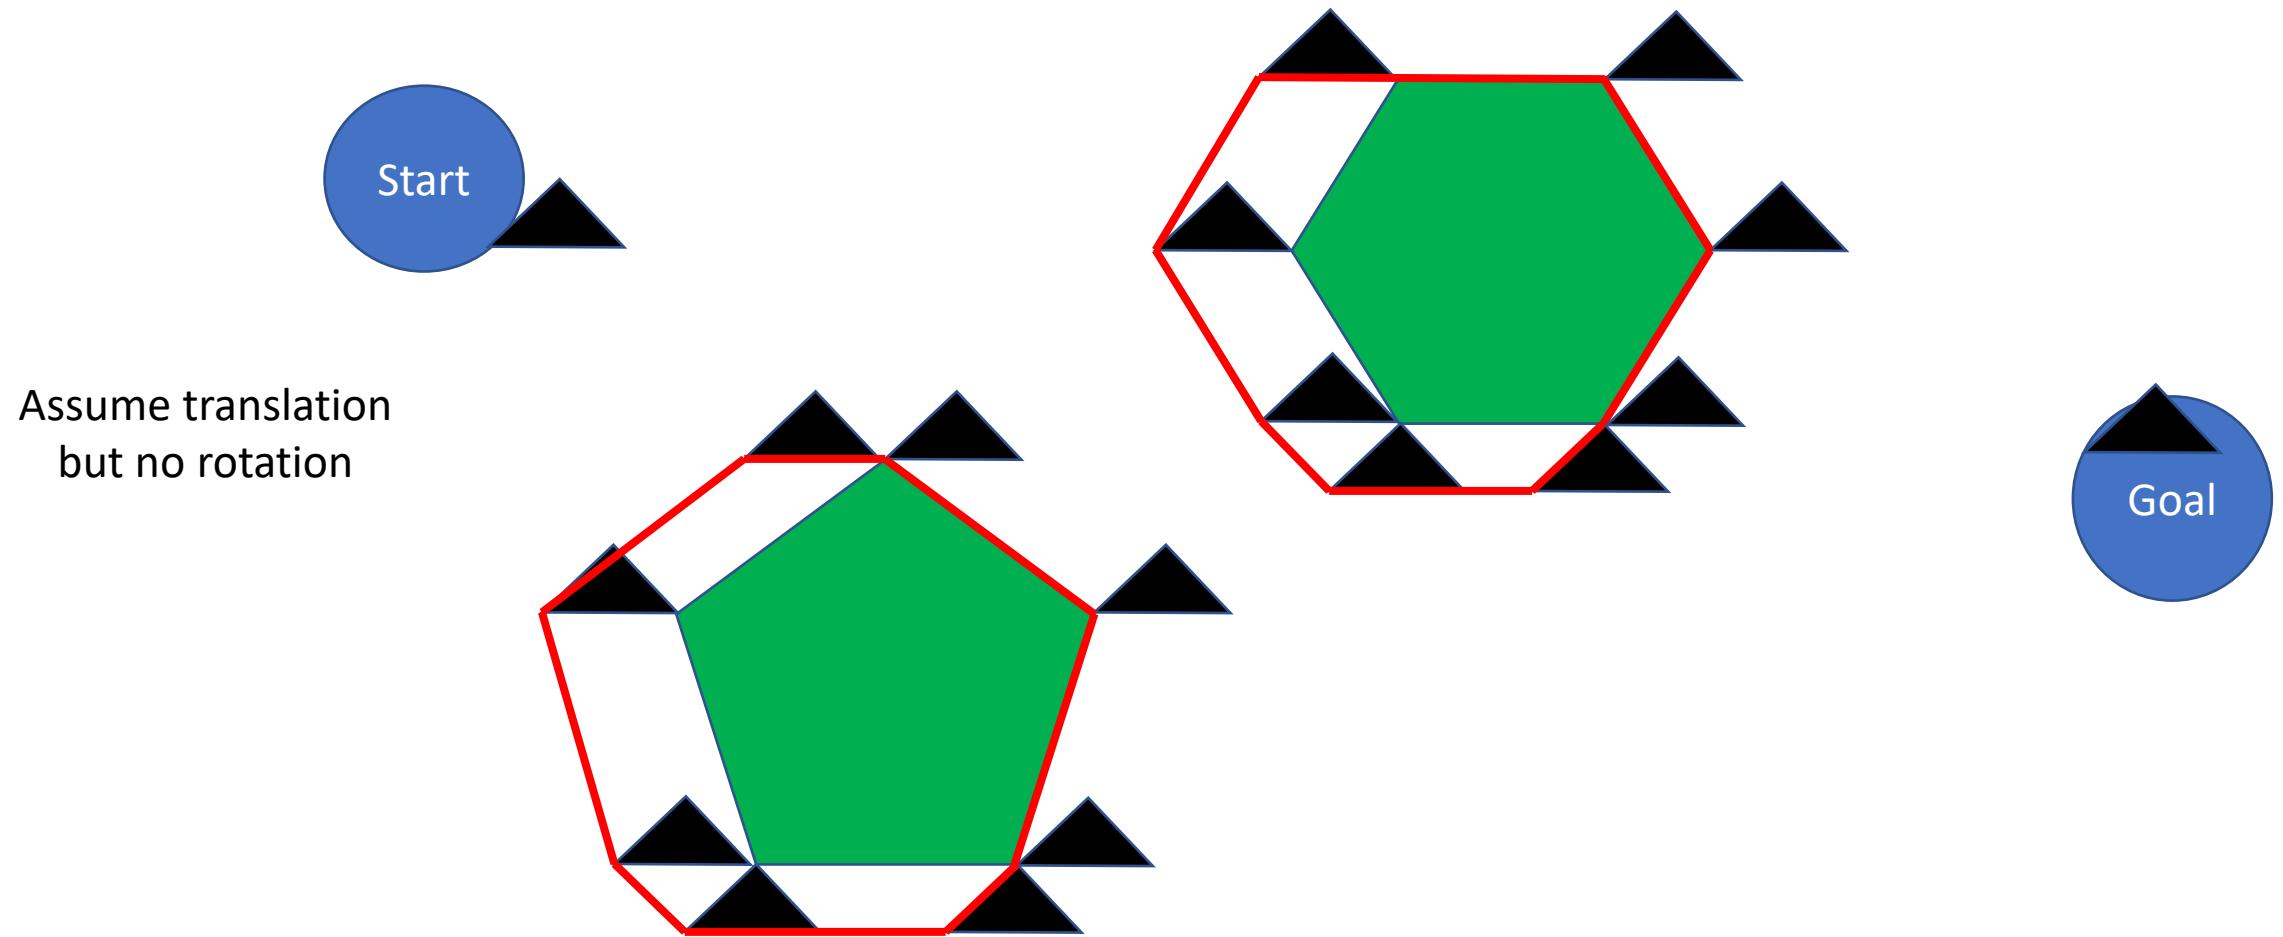
# How do we map path planning to state space search?



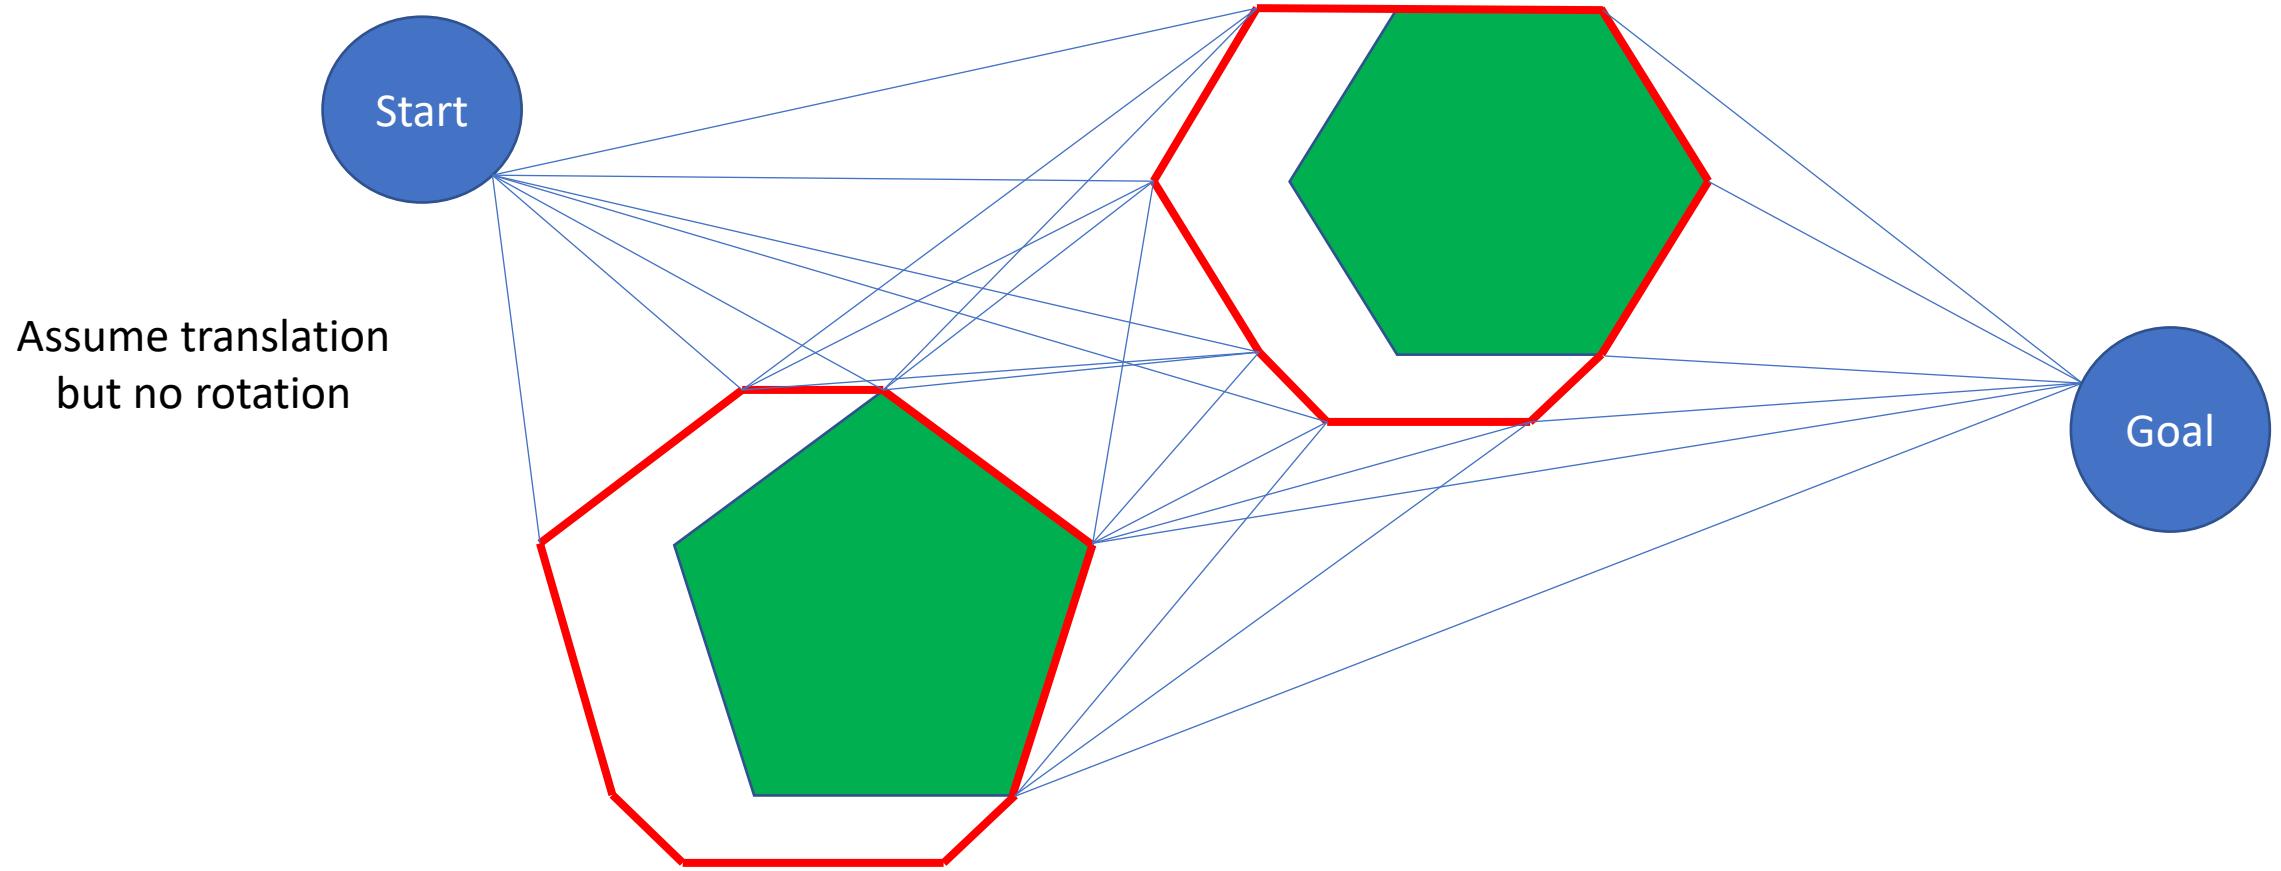
# How do we map path planning to state space search?



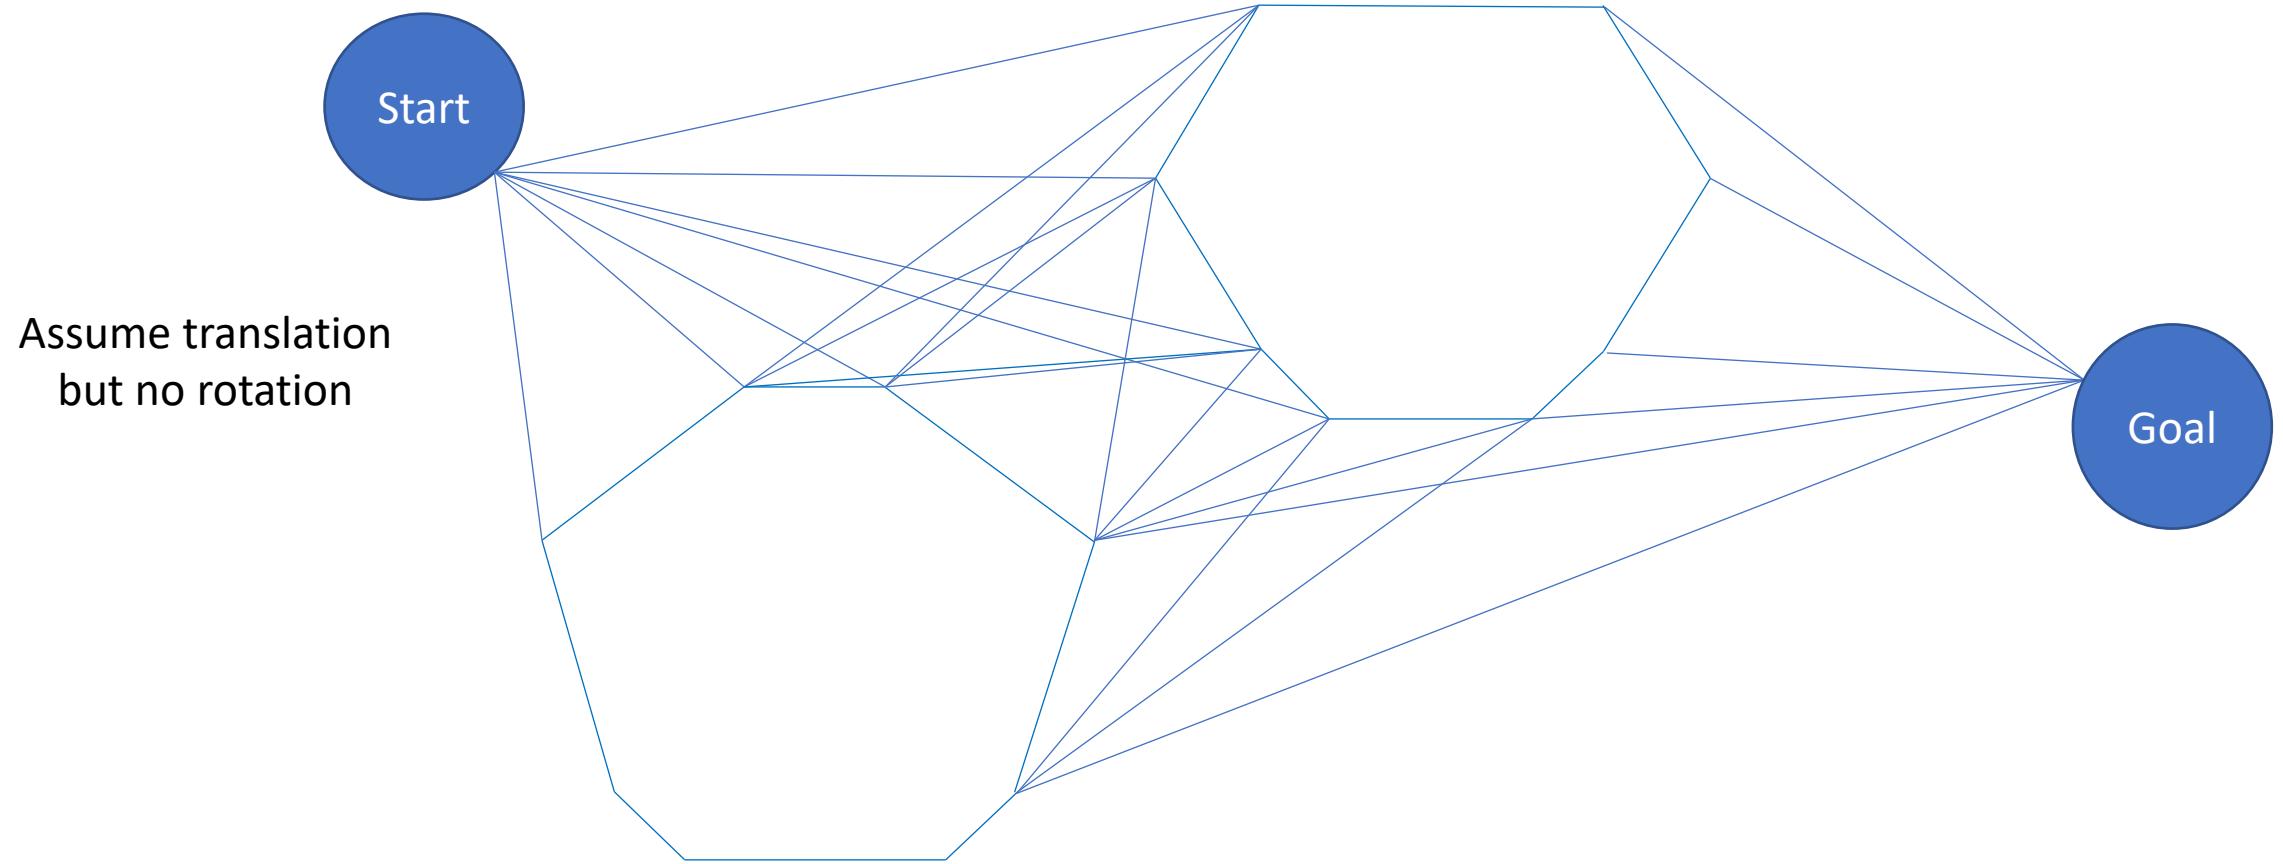
# Create Configuration Space



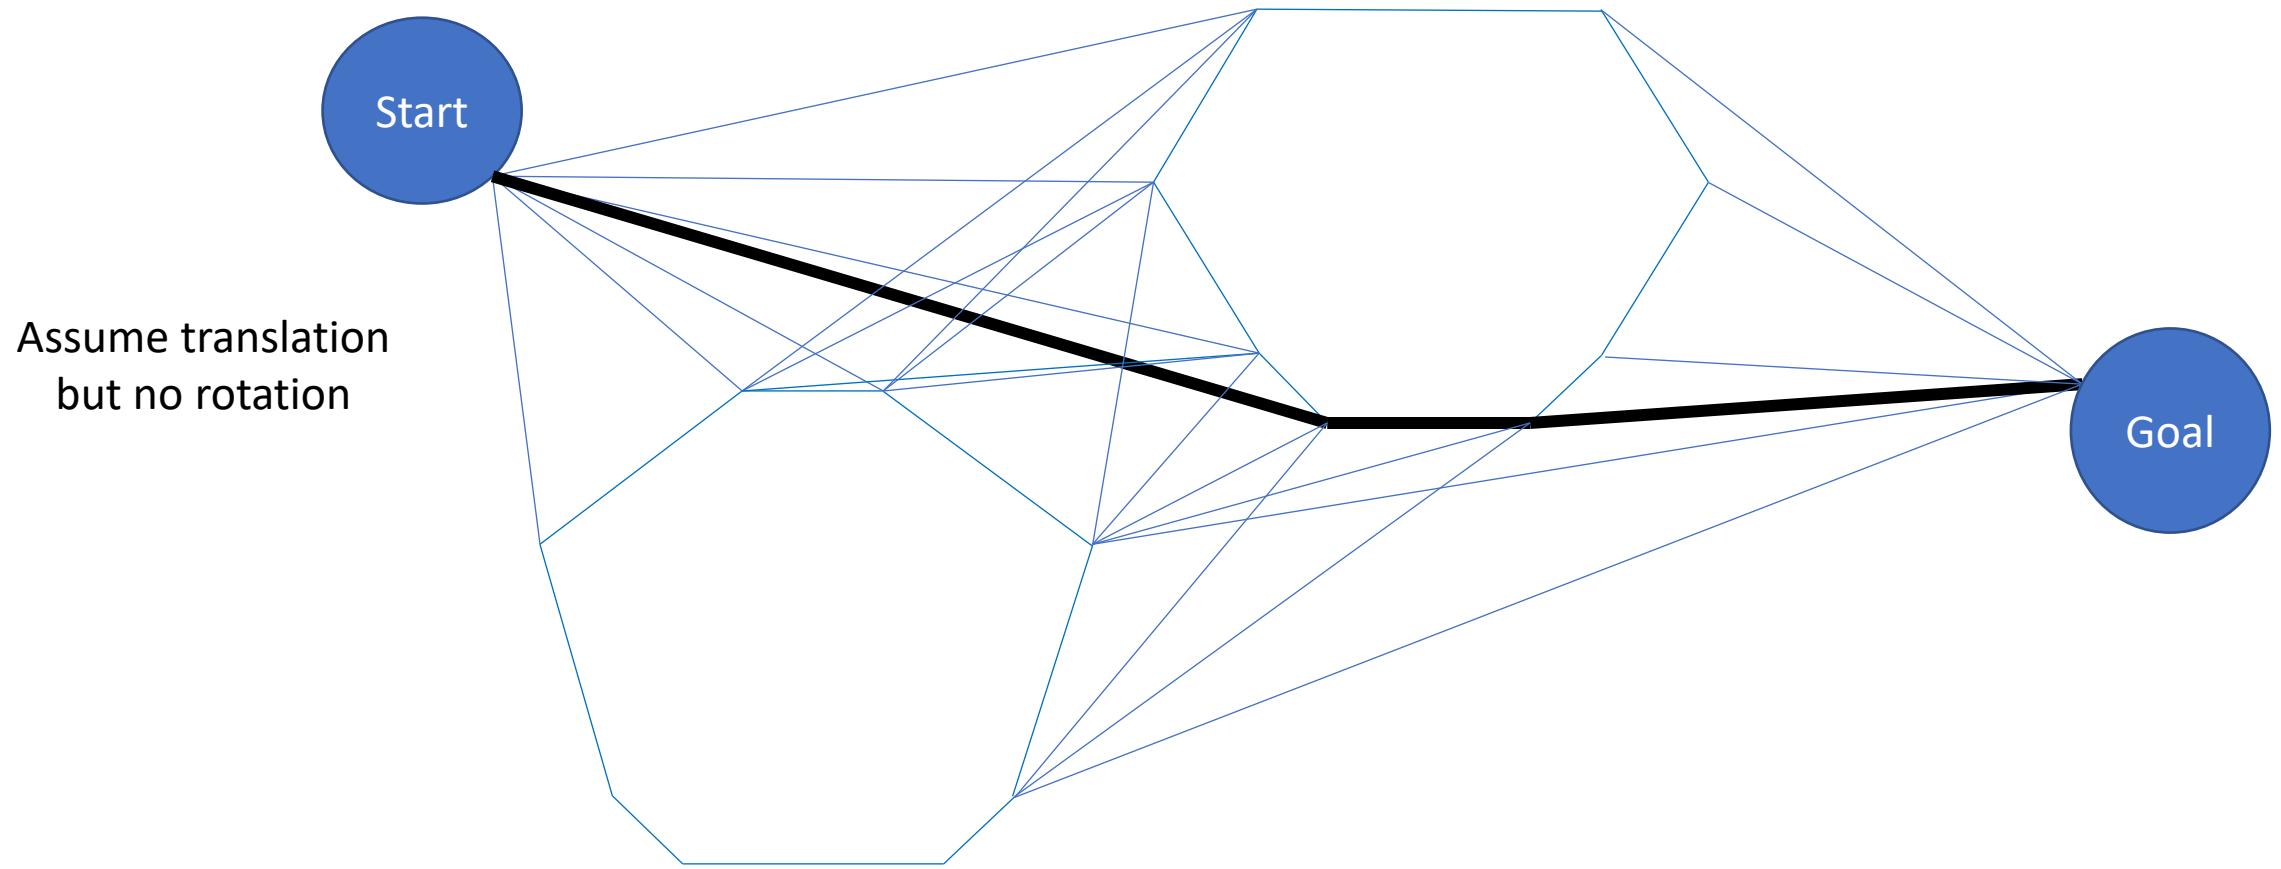
# Map From Continuous Problem to Graph Search: Create Visibility Graph



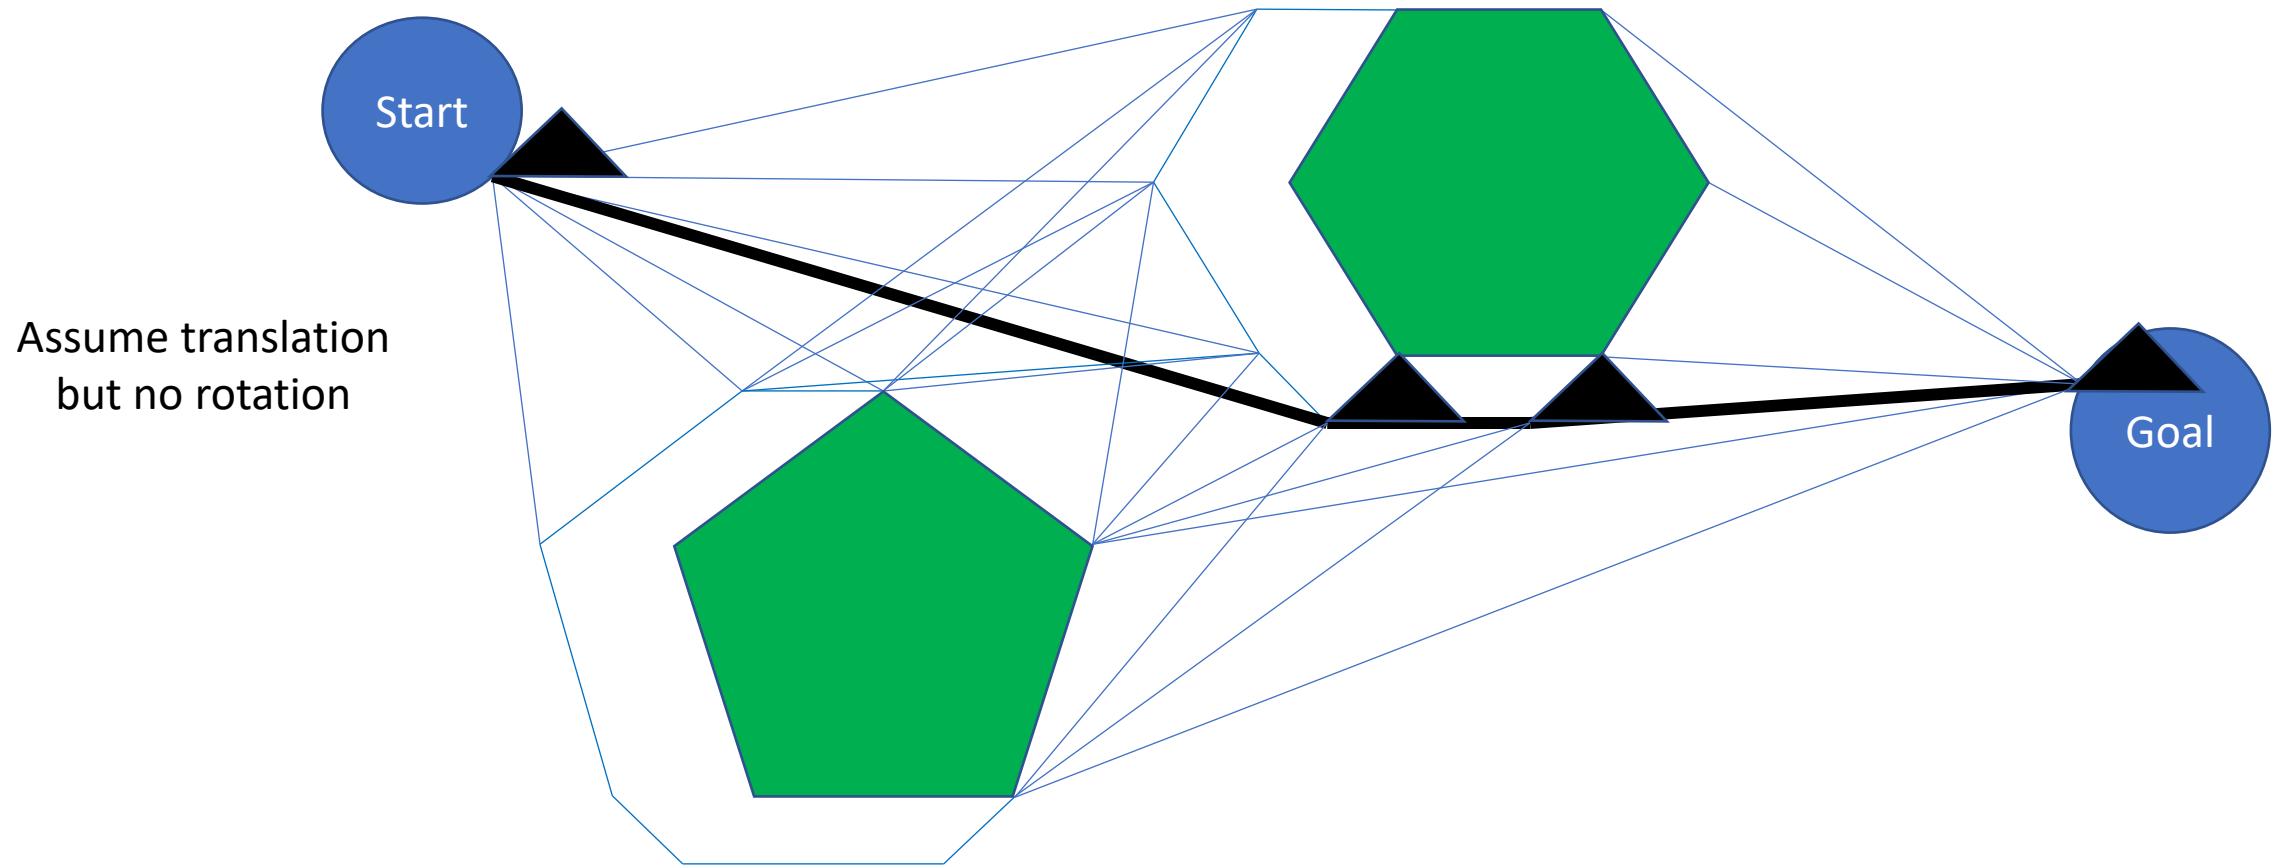
# Map From Continuous Problem to Graph Search: Create Visibility Graph



# Finding Shortest Path (e.g., A\*)

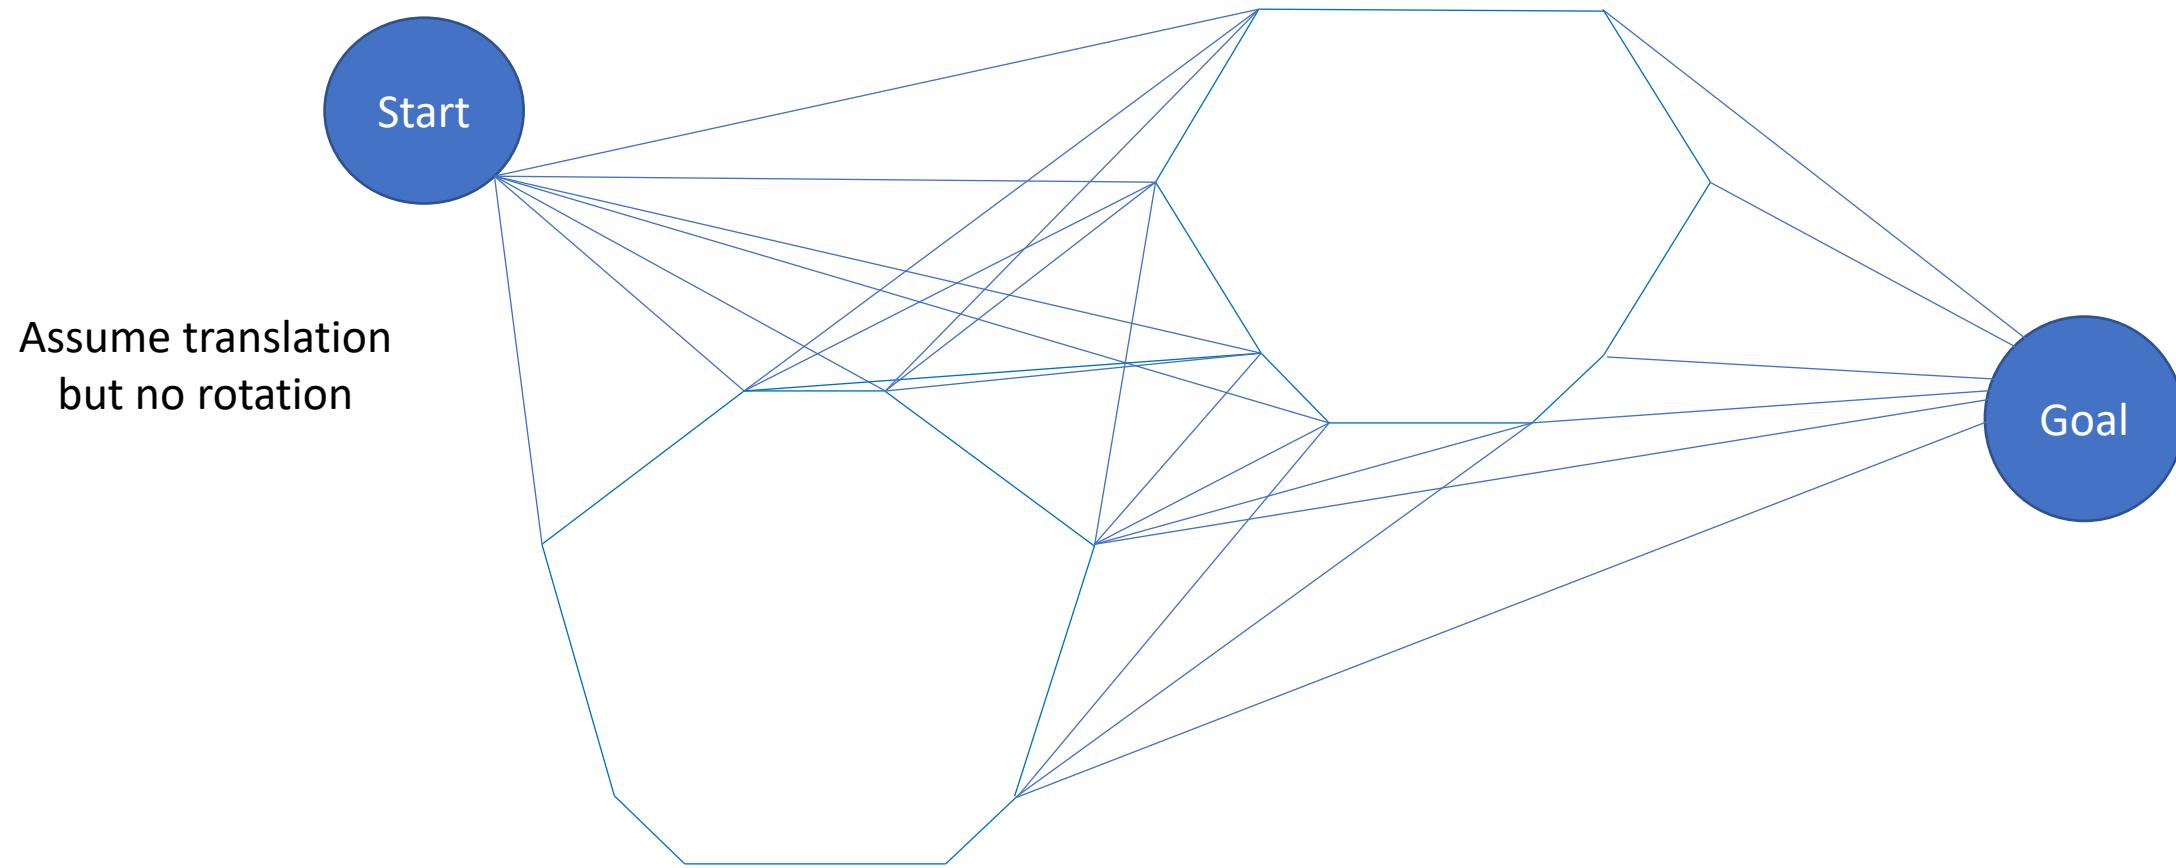


# Finding Shortest Path (e.g., A\*)

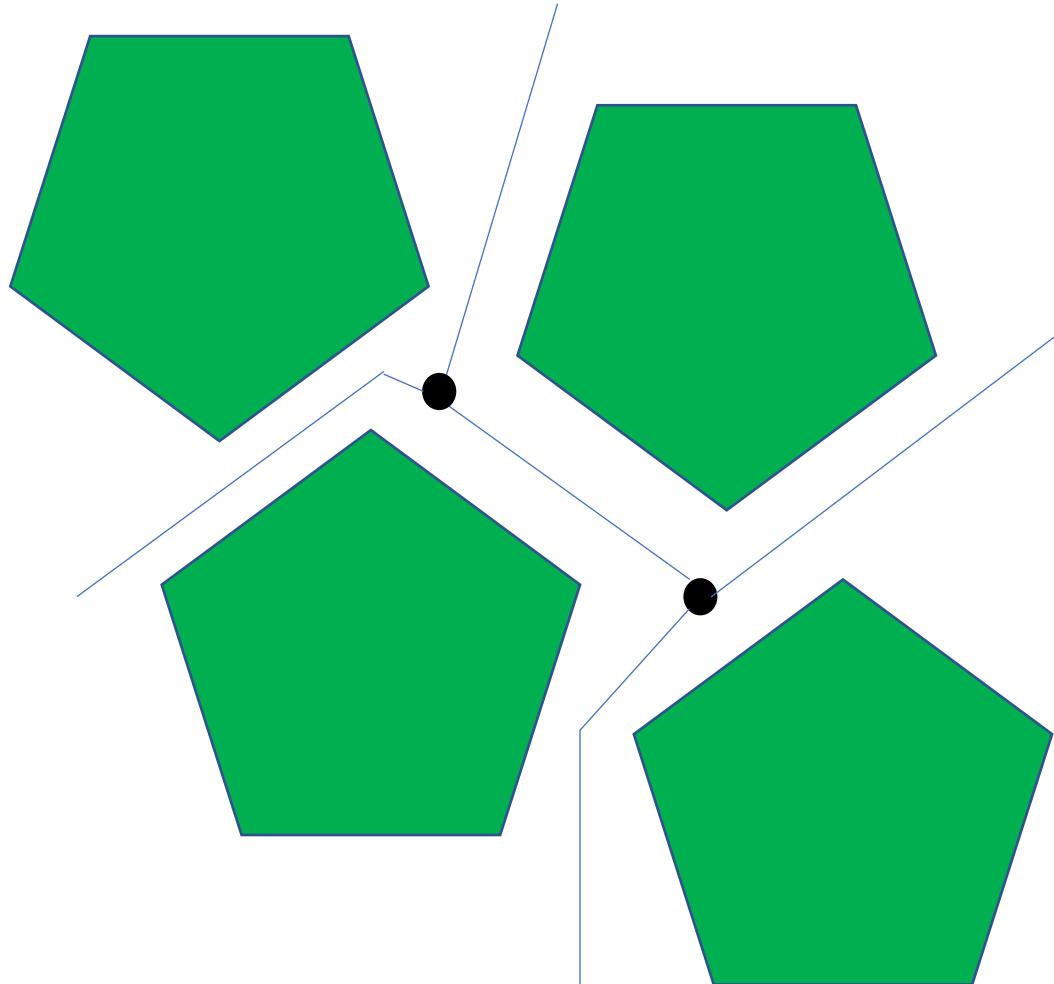


# A Visibility Graph is a Kind of Roadmap

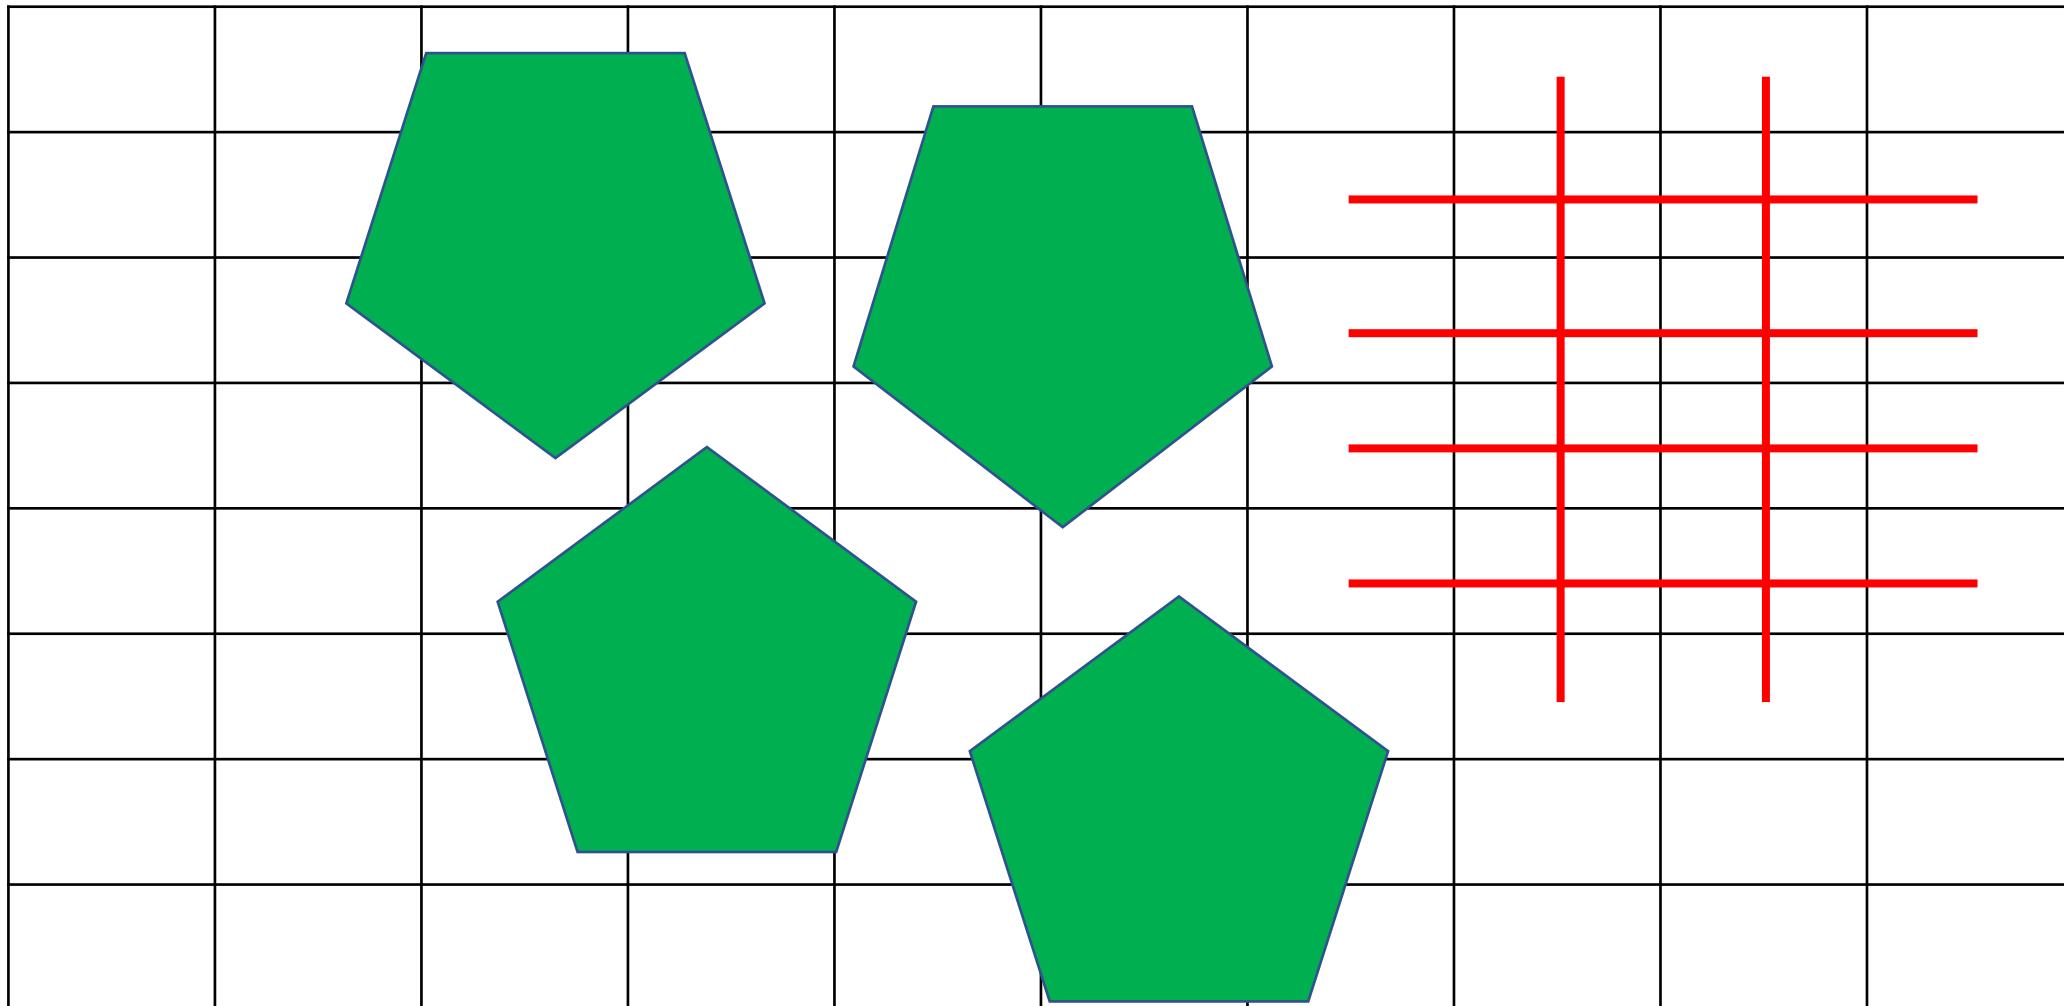
- What are the strengths / weaknesses of roadmaps?
- What are some other types of roadmaps?



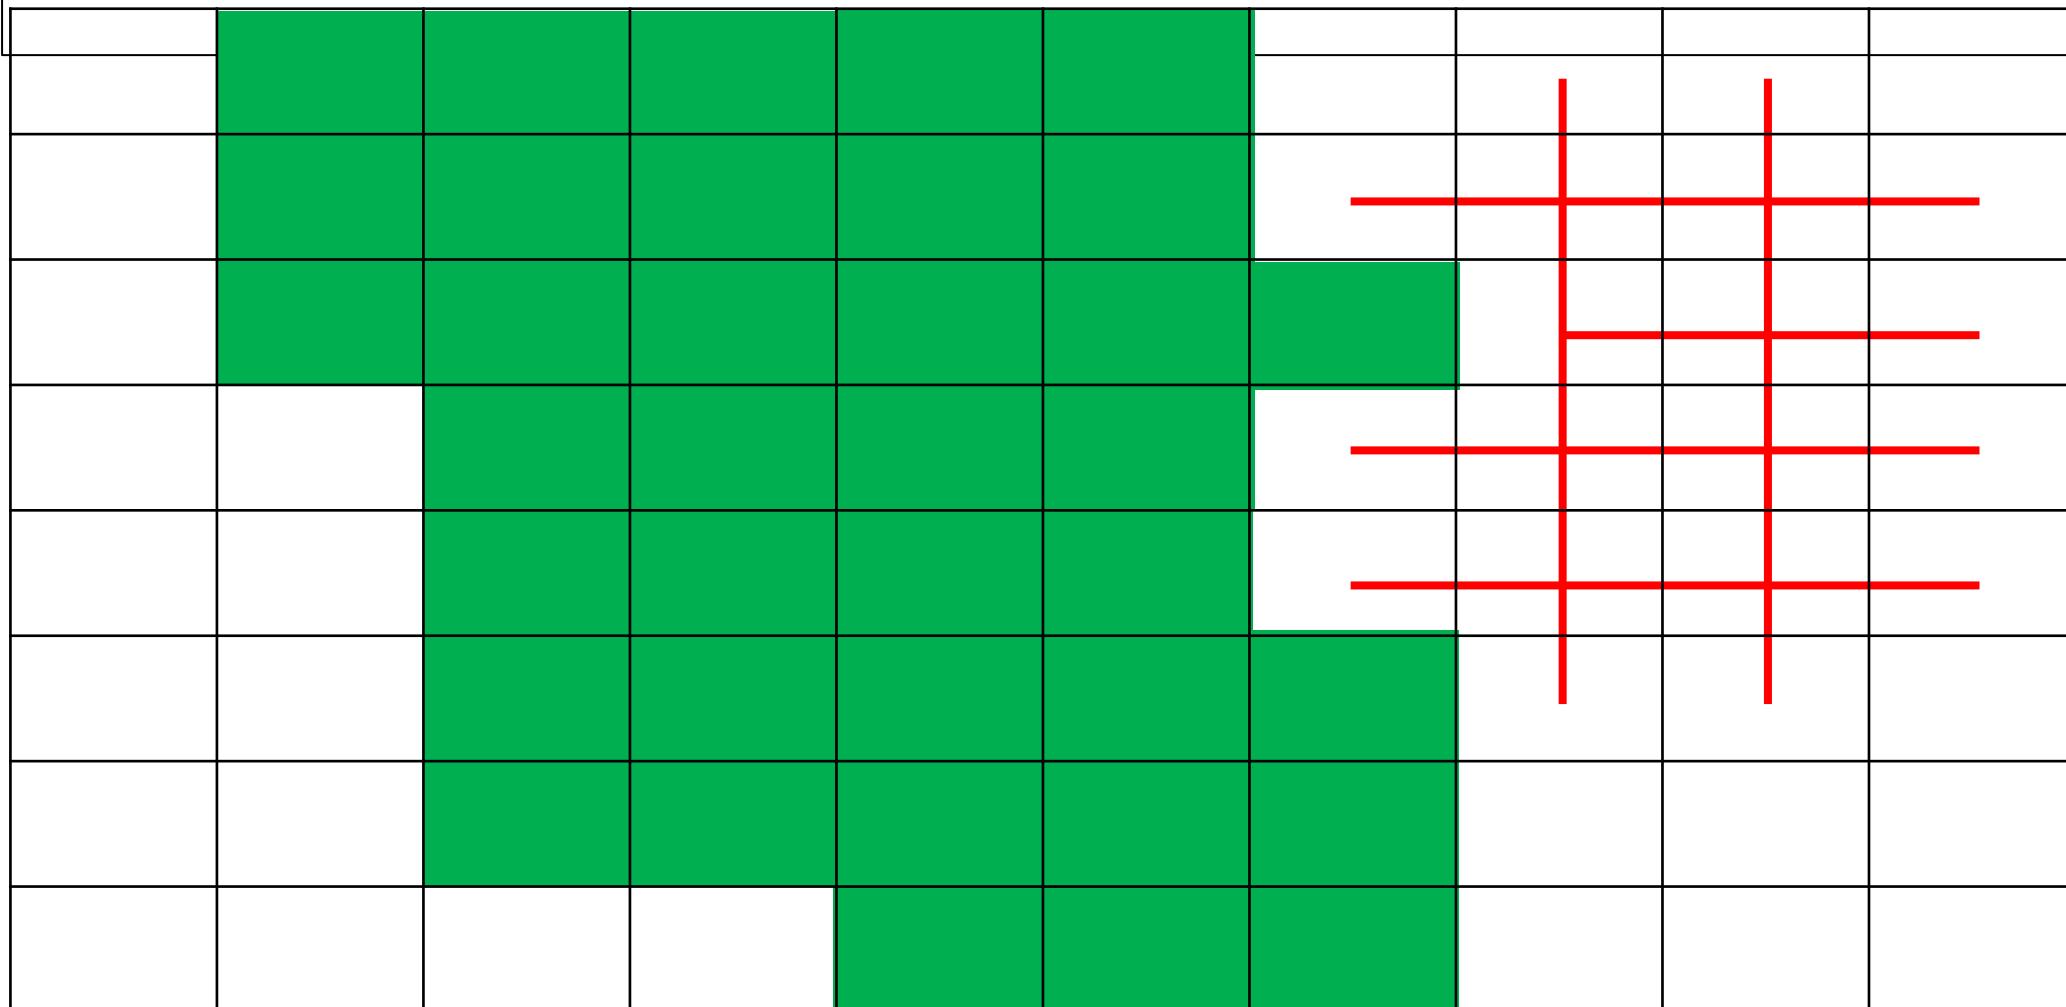
# Voronoi Diagrams



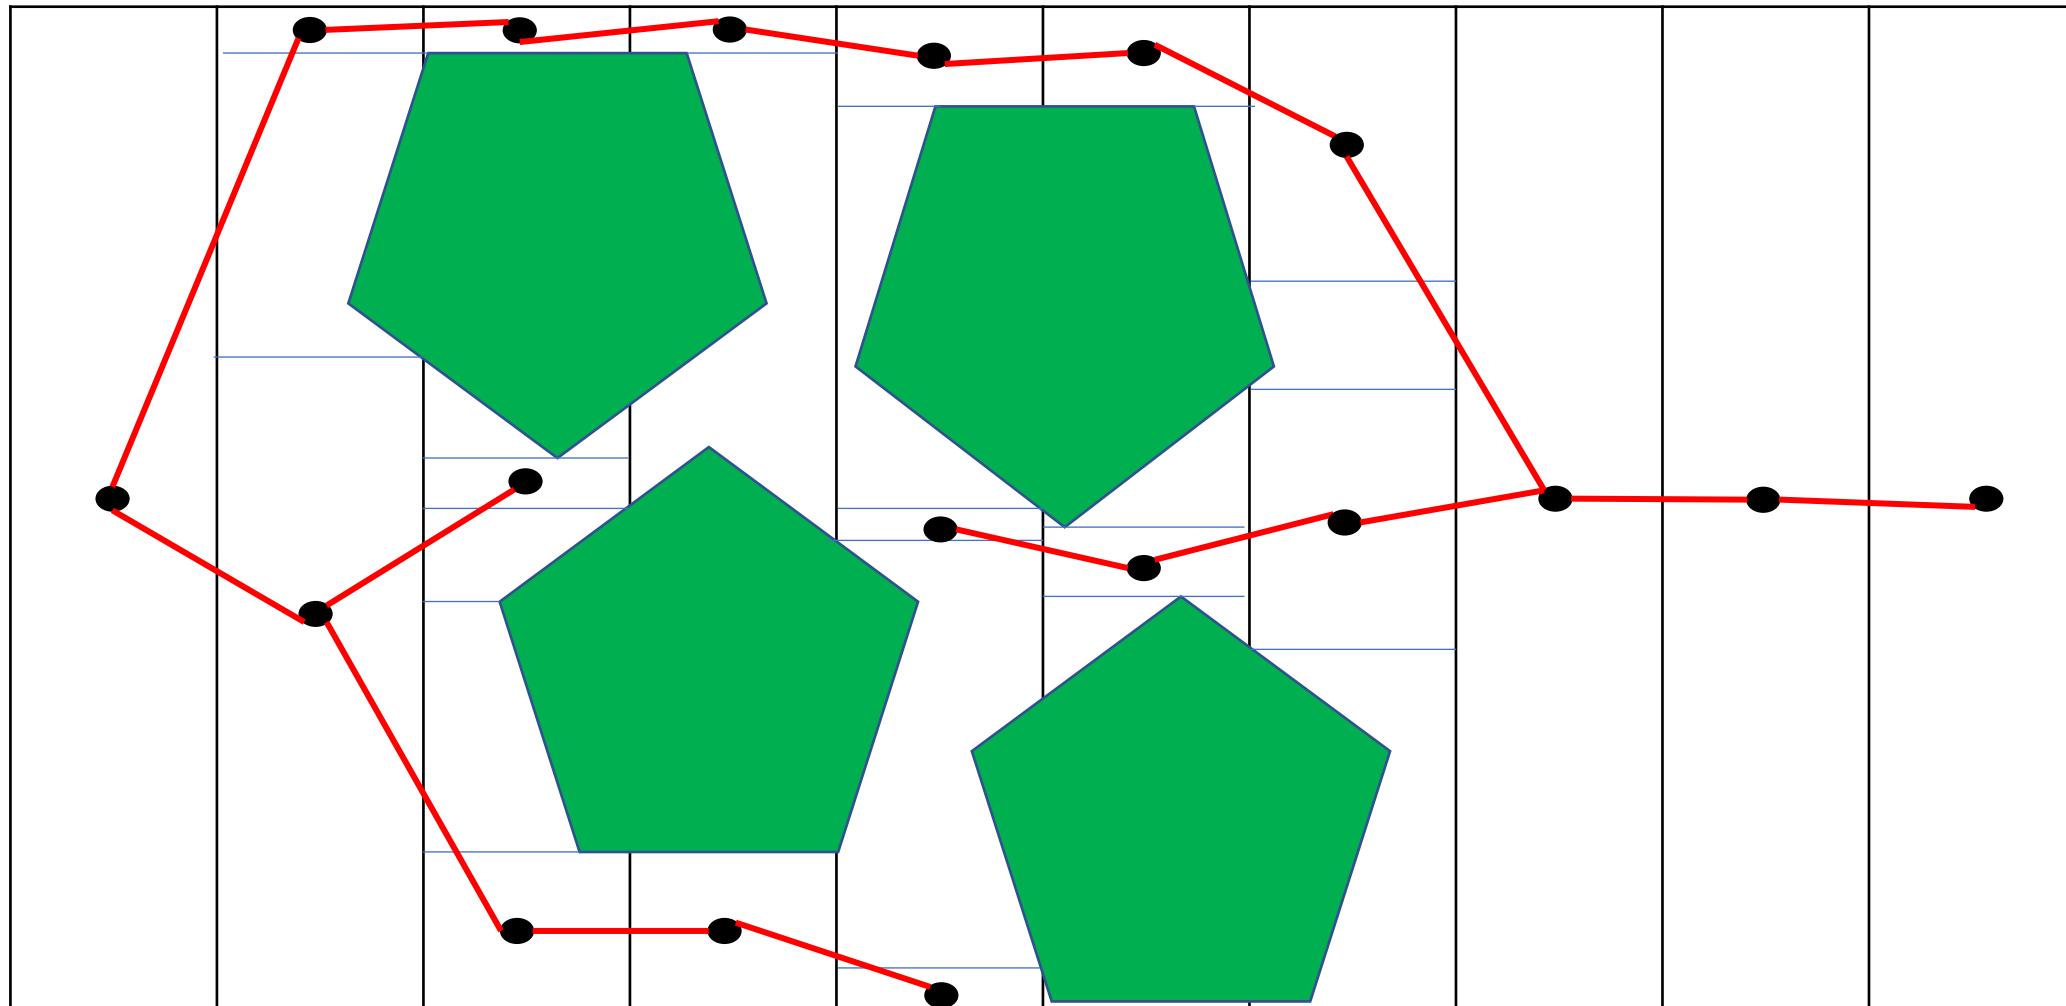
# Roadmaps: Approximate Fixed Cell



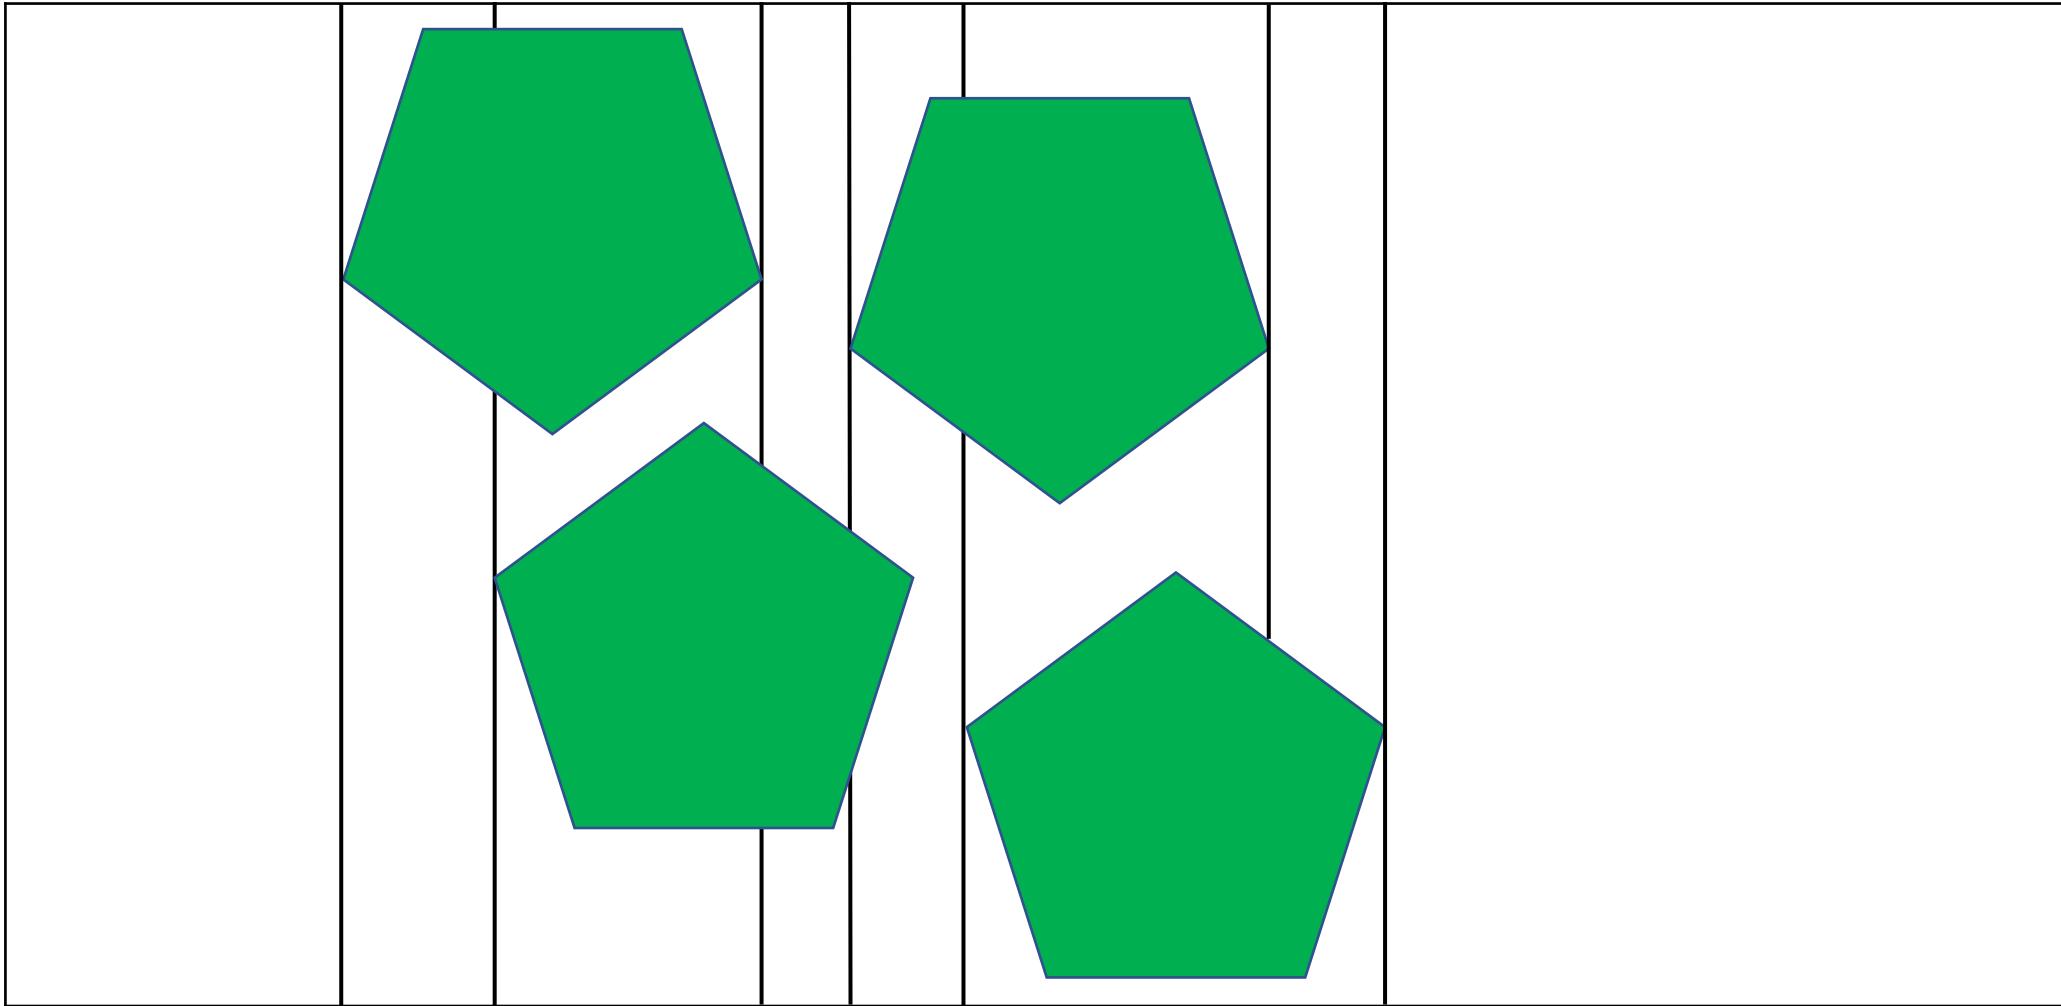
# Roadmaps: Approximate Fixed Cell



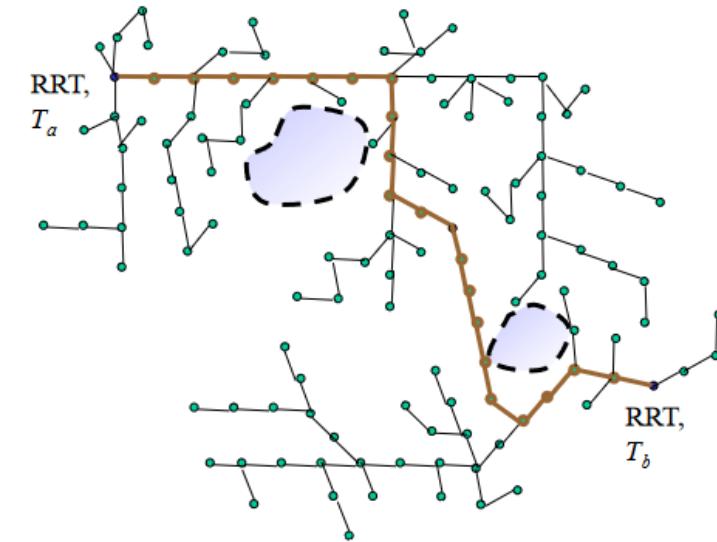
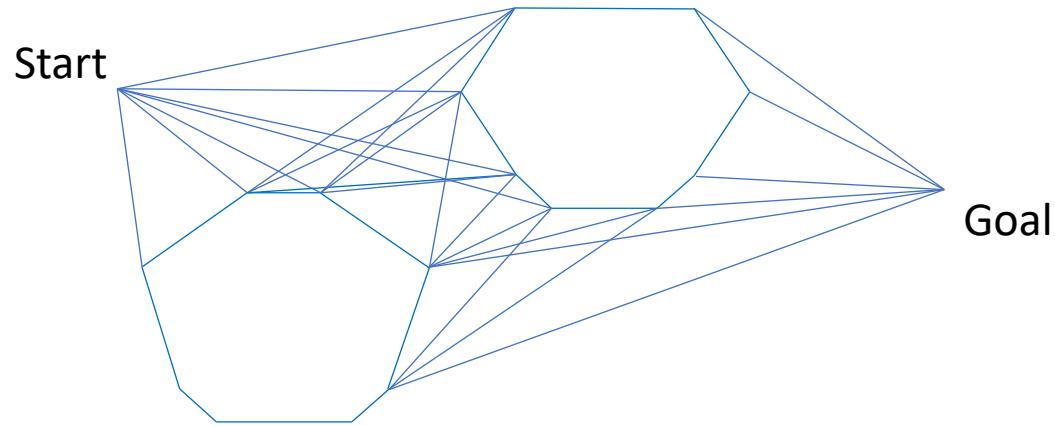
# Roadmaps: Approximate Variable Cell



# Exact Cell Decomposition



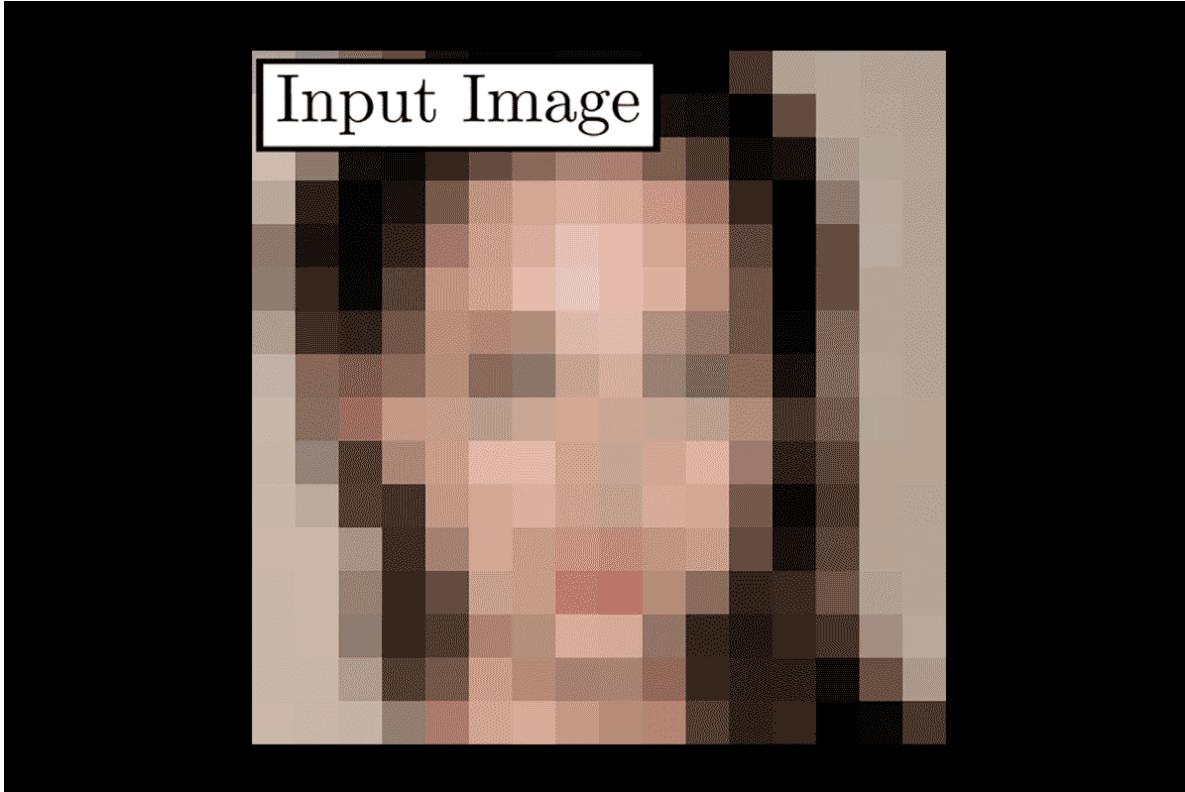
# How do we handle large state spaces?



Rapidly Exploring Random Trees

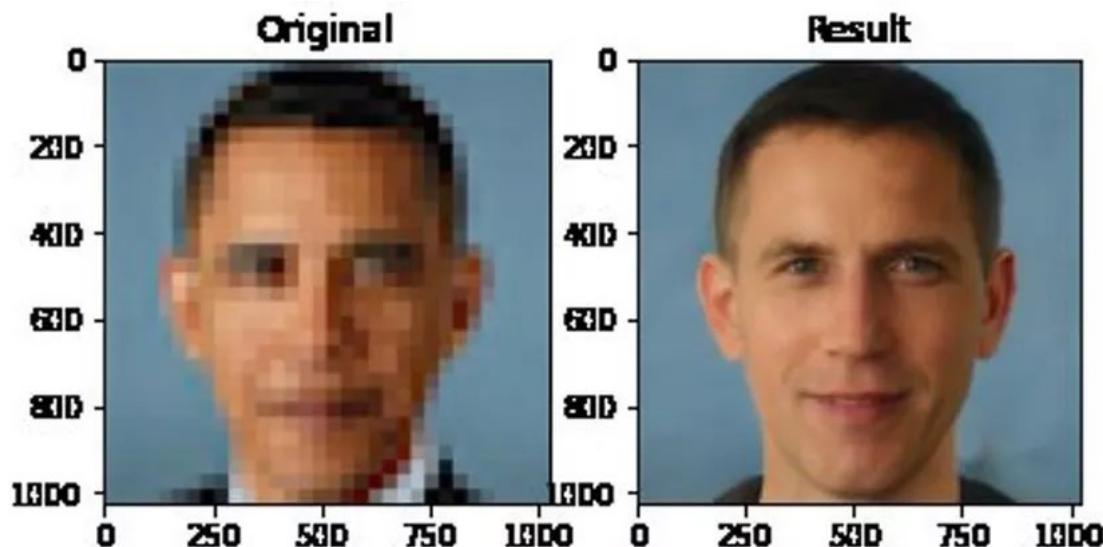
# Mid-lecture Break

# Mid-lecture Break



Menon, S., Damian, A., Hu, S., Ravi, N. and Rudin, C., 2020. PULSE: Self-Supervised Photo Upsampling via Latent Space Exploration of Generative Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 2437-2445).

# Mid-lecture Break



# Mid-lecture Break



Here is my wife @shan\_ness



4:53 PM · Jun 20, 2020 · Twitter Web App

234 Retweets and comments 1.2K Likes



An image of @BarackObama getting upsampled into a white guy is floating around because it illustrates racial bias in #MachineLearning. Just in case you think it isn't real, it is, I got the code working locally. Here is me, and here is @AOC.



4:49 PM · Jun 20, 2020 · Twitter Web App

2.2K Retweets 5.1K Likes



This is @LucyLiu



Lucy Liu

5:06 PM · Jun 20, 2020 · Twitter Web App

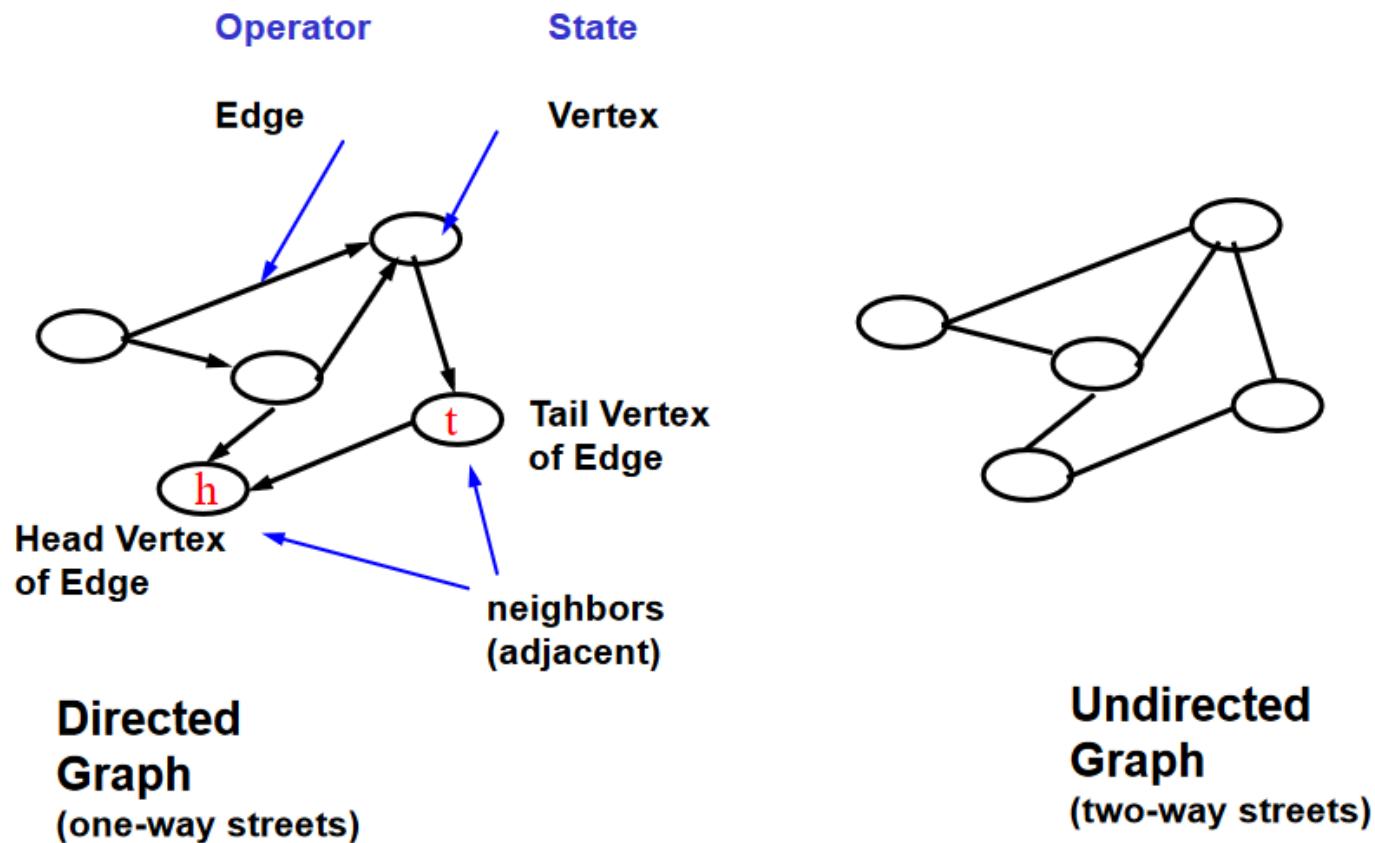
212 Retweets 912 Likes

# Outline:

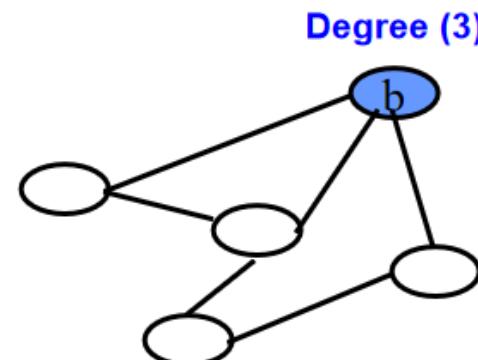
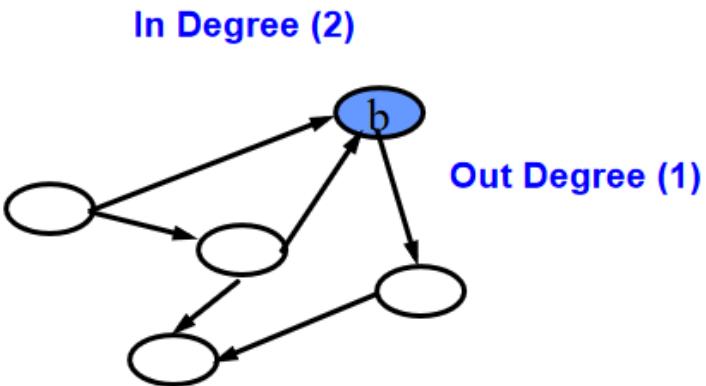
## Problem Solving as State Space Search

- Problem Formulation (Modeling)
  - Problem solving as state space search
- • Formal Representation
  - Graphs and search trees
- Reasoning Algorithms
  - Depth and breadth-first search

# Problem Formulation: A Graph



# Problem Formulation: A Graph



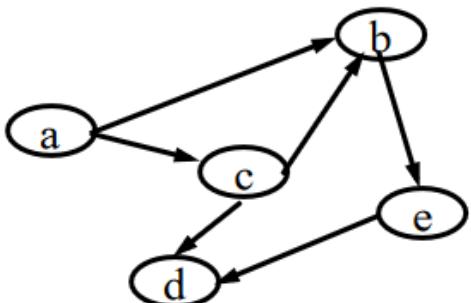
**Directed  
Graph**  
(one-way streets)

**Undirected  
Graph**  
(two-way streets)

# Problem Formulation: A Graph

**Strongly connected graph**

Directed path between all vertices.

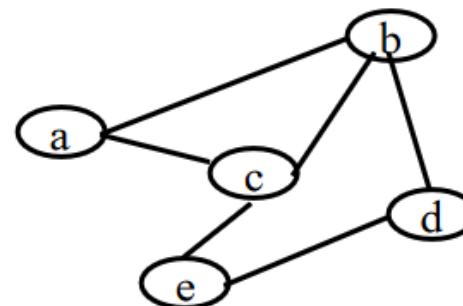


**Connected graph**

Path between all vertices.

**Complete graph**

All vertices are adjacent.



**Sub graph**

Subset of vertices

edges between vertices in Subset

**Directed**

**Graph**

(one-way streets)

**Clique**

A complete subgraph

(All vertices are adjacent).

**Undirected**

**Graph**

(two-way streets)

# Specifying a Graph

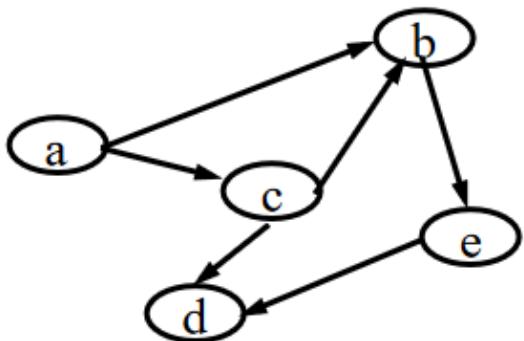
Notation:

$\langle a, b, \dots, n \rangle$

$\{a, b, \dots, n\}$

an ordered list of elements, a, b, ...

an unordered set of distinguished elements

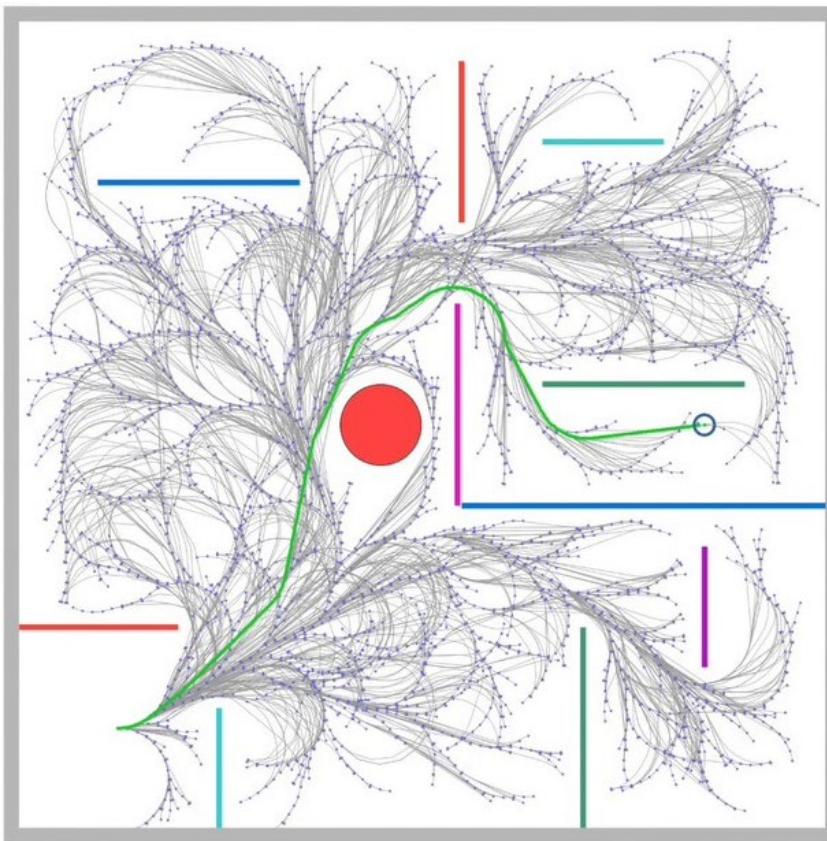


Vertices  $V = \{a, b, c, d, e\}$

Edges  $E = \{\langle a, b \rangle, \langle a, c \rangle, \langle b, e \rangle, \langle c, b \rangle, \langle c, d \rangle, \langle e, d \rangle\}$

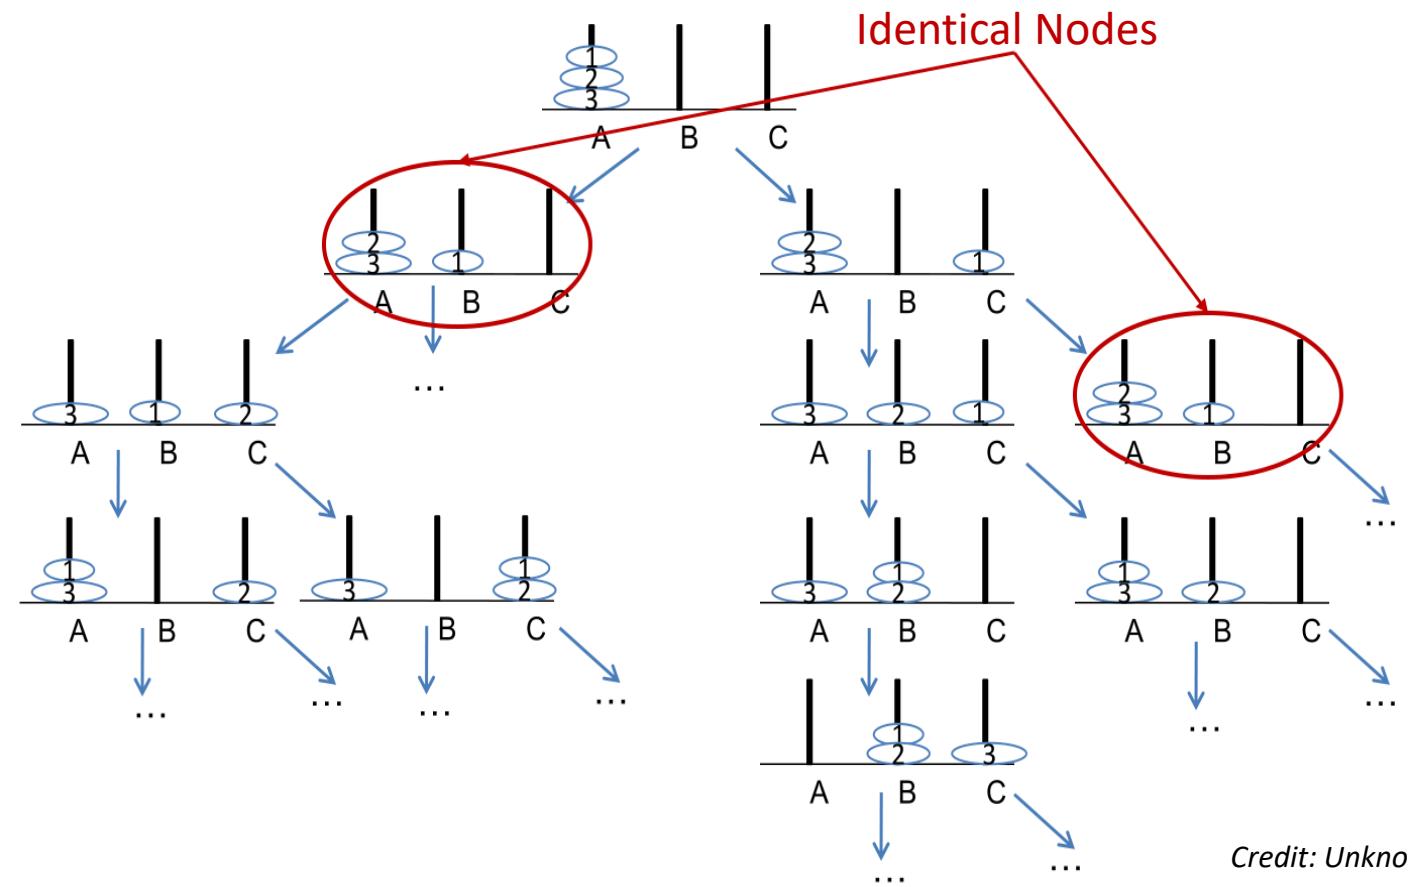
# Examples of Graphs

Roadmap

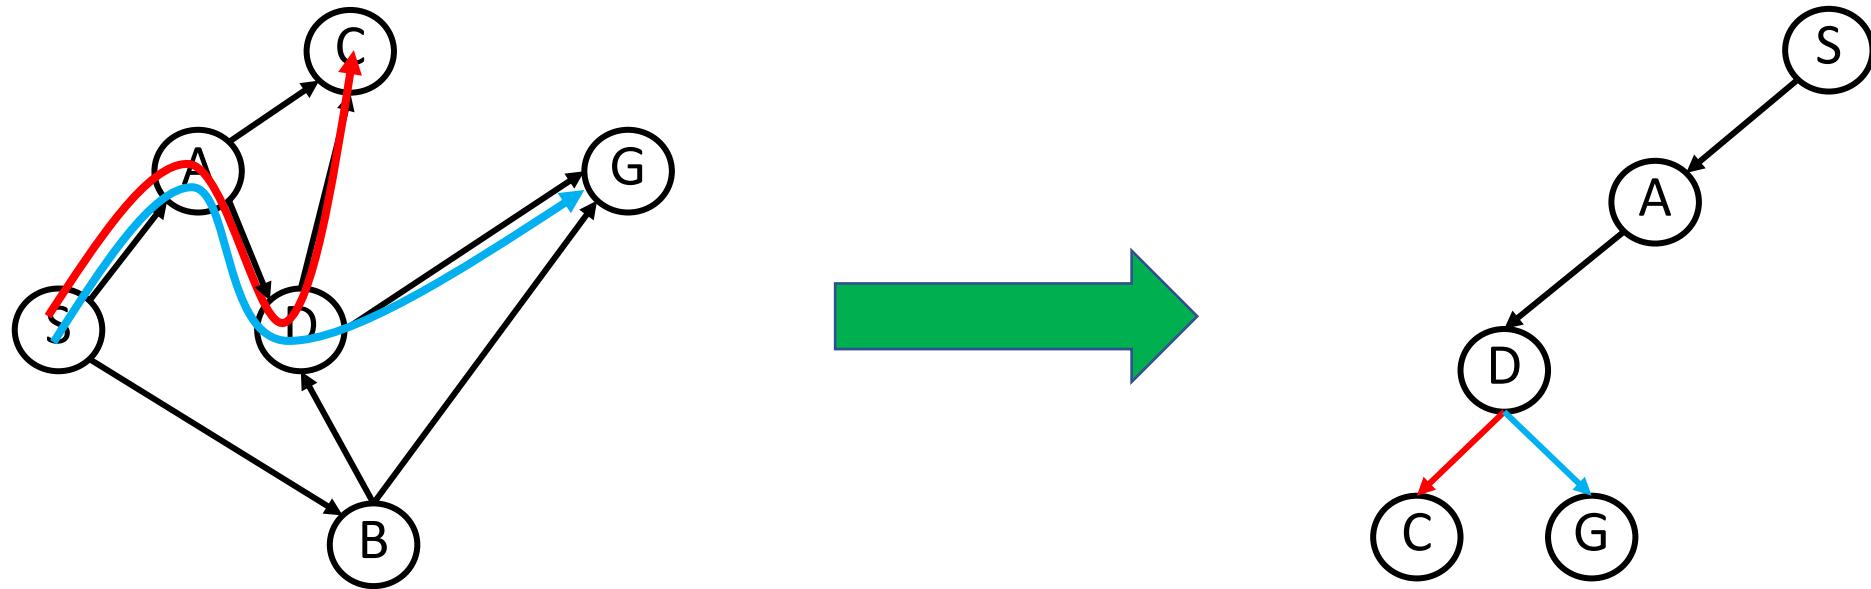


Credit: Olzhas Adiyatov and H. Atakan Varol

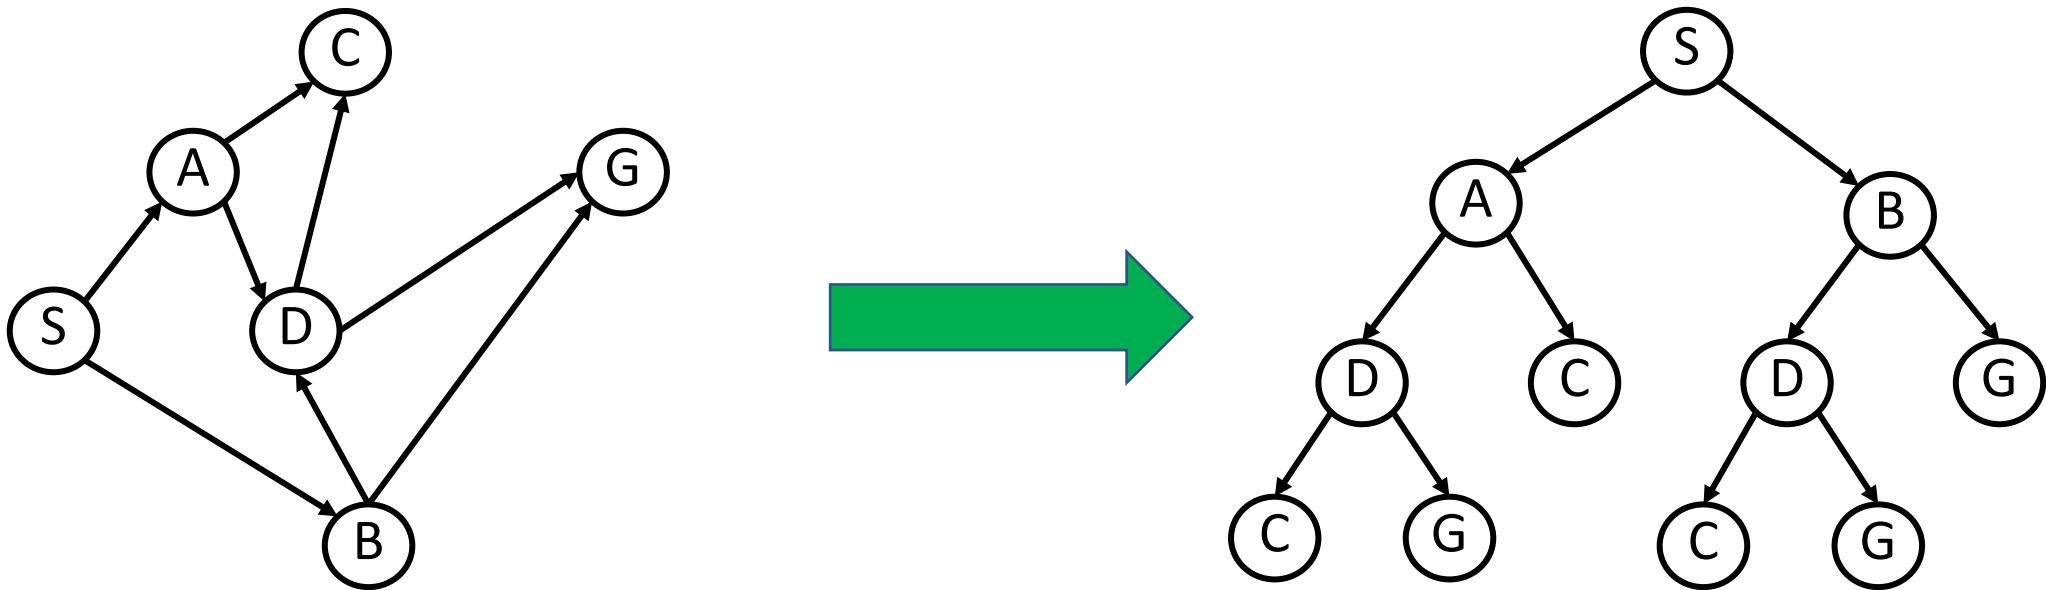
Planning Actions  
(Graph of possible states of the world)



# A Problem Solver Searches Through All Simple Paths



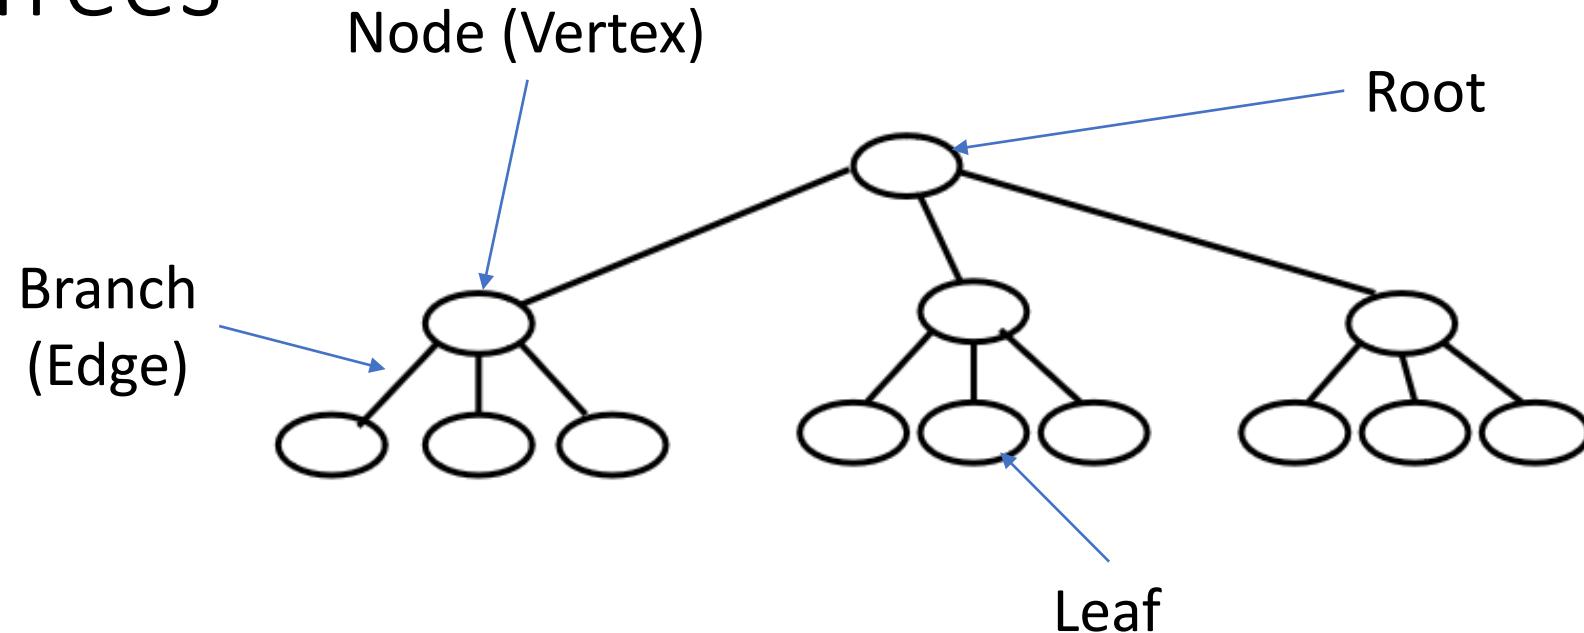
# Search Tree Denotes All Simple Paths



Enumeration is:

- Complete
- Sound

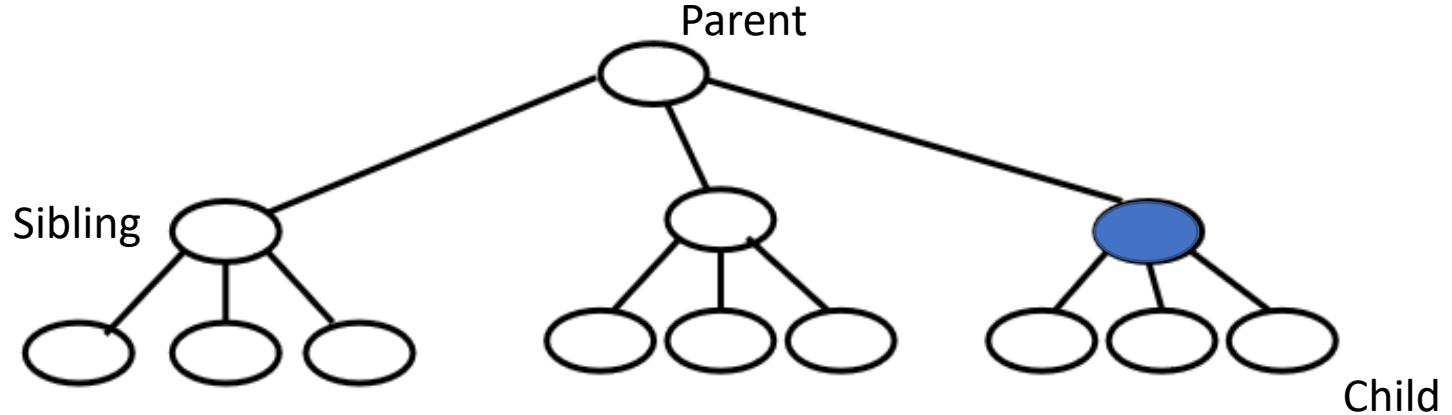
# Search Trees



A tree  $T$  is a directed graph such that:

- There exists exactly one undirected path between any pair of vertices
- *In degree* of each vertex is 1

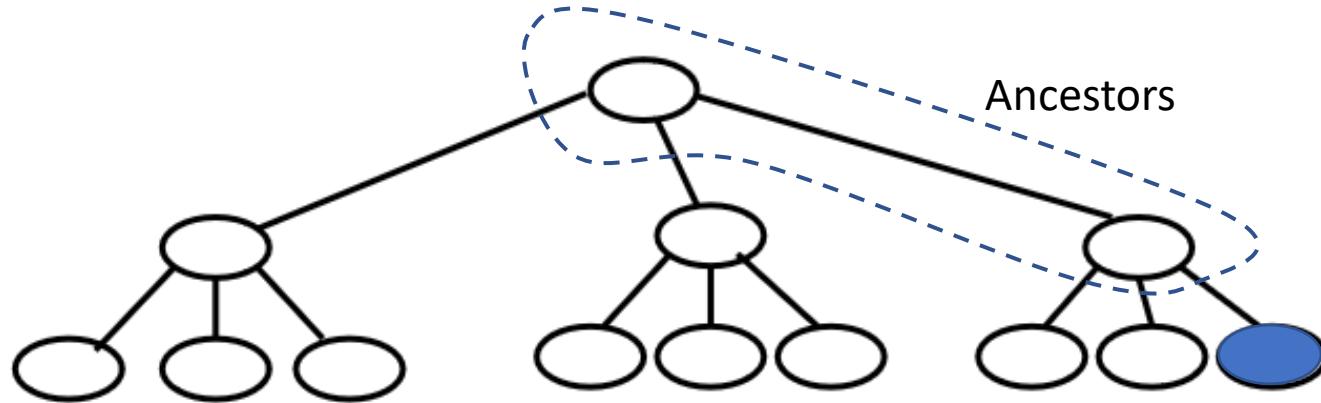
# Search Trees



A tree  $T$  is a directed graph such that:

- There exists exactly one undirected path between any pair of vertices
- *In degree* of each vertex is 1

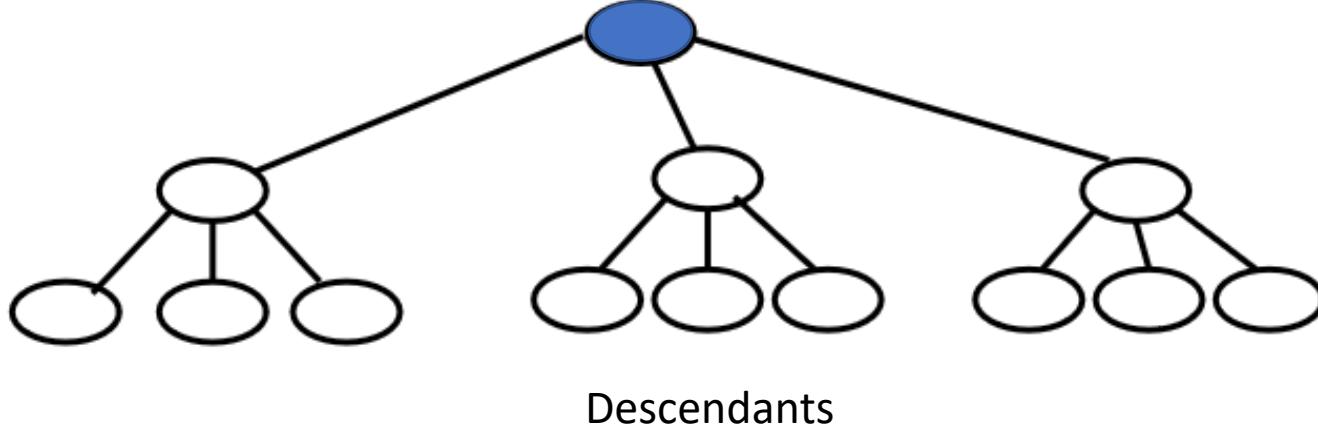
# Search Trees



A tree  $T$  is a directed graph such that:

- There exists exactly one undirected path between any pair of vertices
- *In degree* of each vertex is 1

# Search Trees



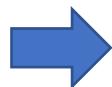
A tree  $T$  is a directed graph such that:

- There exists exactly one undirected path between any pair of vertices
- *In degree* of each vertex is 1

# Outline:

## Problem Solving as State Space Search

- Problem Formulation (Modeling)
  - Problem solving as state space search
- Formal Representation
  - Graphs and search trees
- Reasoning Algorithms
  - Depth and breadth-first search



# Taxonomy of Search

Class	Types	Description
<b>Blind</b> (uninformed)	Depth-First Breadth-First Iterative-Deepening	Systematic exploration of the whole tree until the goal is found
<b>Optimal</b> (informed)	A* Branch & Bound	Use path “length” measure to find the shortest path
<b>Heuristic</b> (informed)	Hill-Climbing Best-First Beam	Use a heuristic measure of goodness of a node

# Simple Search Algorithm

Let Q be a list of partial paths,

S be the start node and

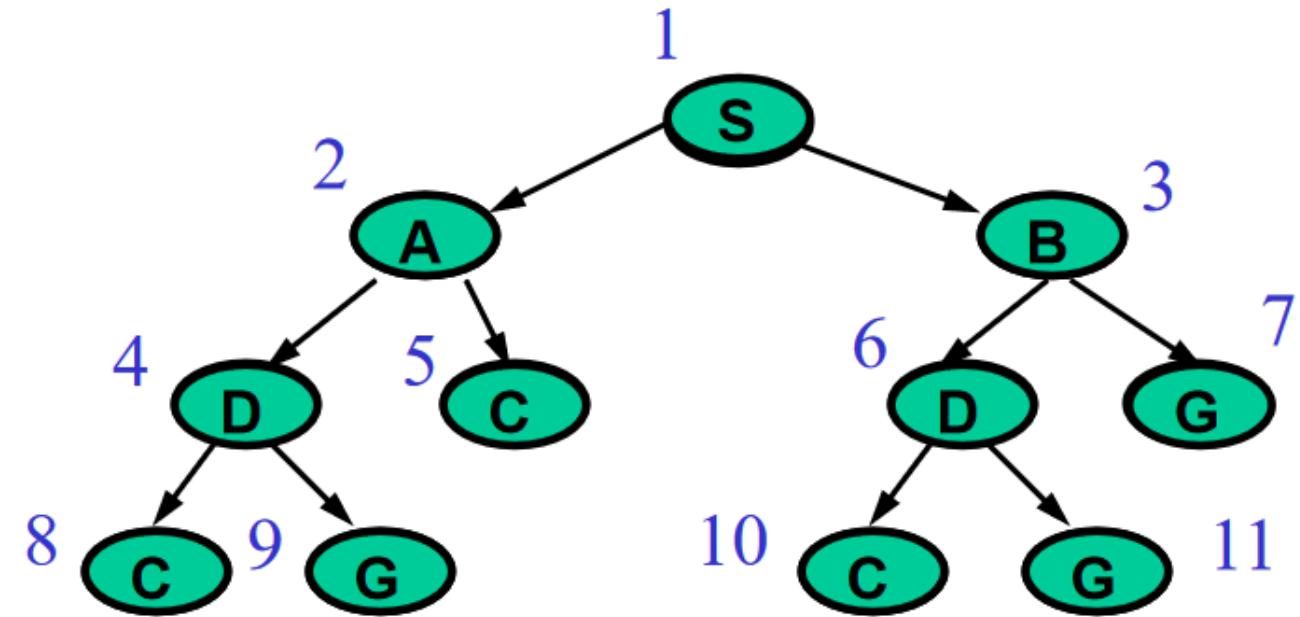
G be the goal node.

1. Initialize Q with a partial path  $\langle S \rangle$
2. If Q is empty, fail. Else, pick a partial path N from Q
3. If  $\text{head}(N) = G$ , return N // Goal reach!
4. Else:
  - a) Remove N from Q
  - b) Find all children of  $\text{head}(N)$  and create a one-step extension of N to each child
  - c) Add all extended paths to Q // How we add to Q defines search type
  - d) GOTO step 2

# Breadth-First Search

Idea: After visiting a node:

- Visit each child
- Visit each child's sibling before the child's children

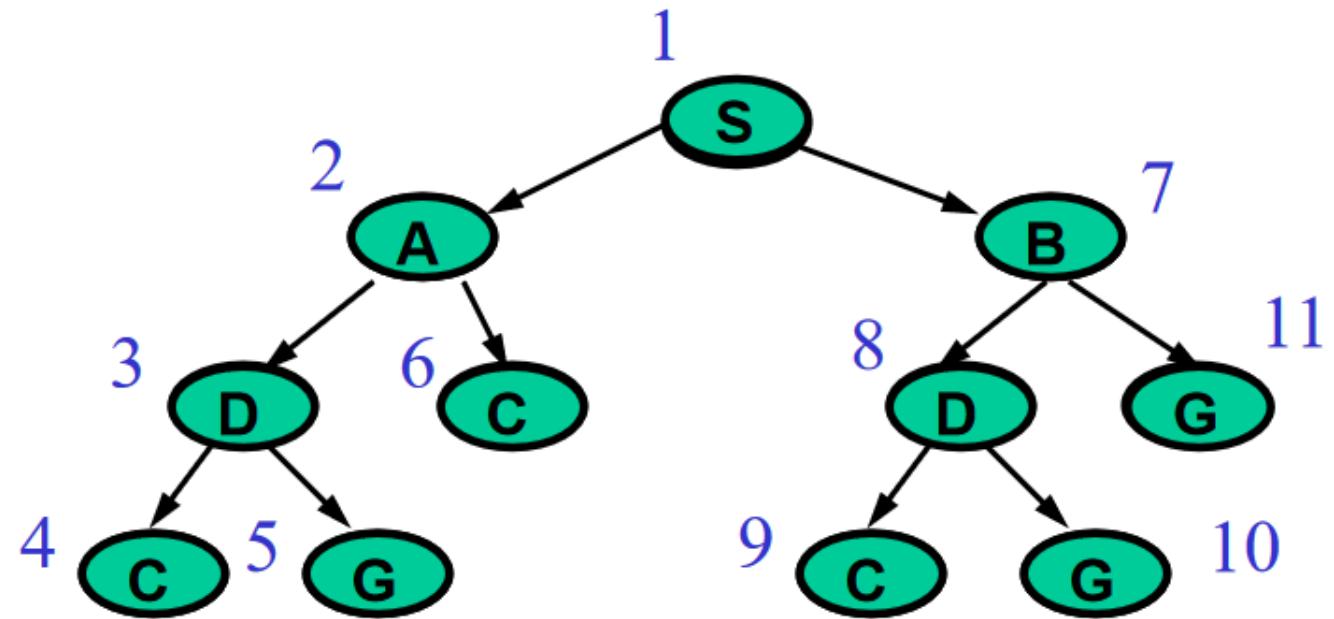


What kind of queue is this?  
→ First-in-first-out (FIFO)

# Depth-First Search

Idea: After visiting a node:

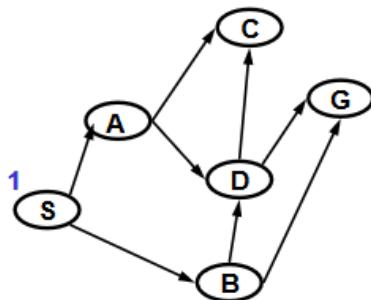
- Visit each child
- Visit each child's children before the child's sibling



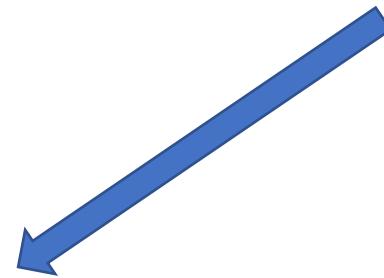
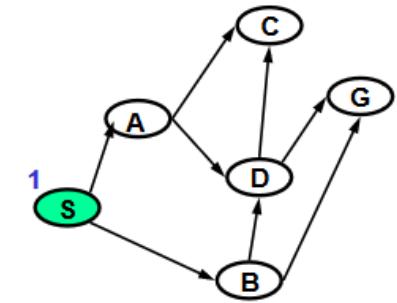
What kind of queue is this?  
→ Last-in-first-out (LIFO)

# DFS Example

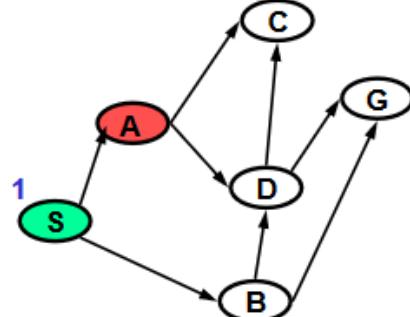
	Q
1	(S)
2	
3	
4	
5	
6	



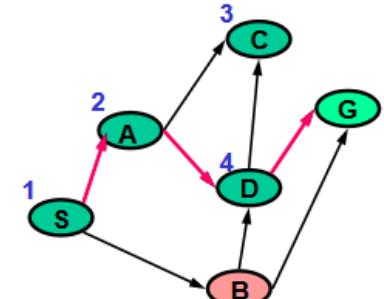
	Q
1	(S)
2	
3	
4	
5	



	Q
1	(S)
2	(A S)
3	
4	
5	

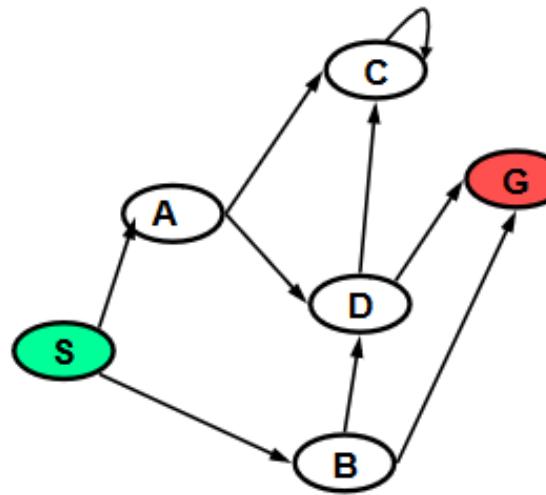


	Q
1	(S)
2	(A S) (B S)
3	(C A S) (D A S) (B S)
4	(D A S) (B S)
5	(C D A S) (G D A S) (B S)
6	(G D A S) (B S)



# Issue

With DFS, will we ever reach G?



Is effort wasted even when you don't have cycles?

# Solution: Visited List

Idea:

- Keep track of nodes already visited
- Do not place expanded path on Q if head is a visited node

Does this maintain correctness?

- Any goal reachable from a node that was visited a second time would be reachable from that node the first time

Does this always improve efficiency?

- Visits only a subset of the original path

# Simple Search Algorithm (Handling Cycles)

Let Q be a list of partial paths,

S be the start node and

G be the goal node.

1. Initialize Q with a partial path  $\langle S \rangle$ ; set Visited = {}
2. If Q is empty, fail. Else, pick a partial path N from Q
3. If  $\text{head}(N) = G$ , return N //Goal reach!
4. Else:
  - a) Remove N from Q
  - b) Find all children of  $\text{head}(N)$  not in Visited and create a one-step extension of N to each child
  - c) Add all extended paths to Q
  - d) Add Children of head(N) to Visited**
  - e) GOTO step 2

# Elements of Algorithmic Design

## Algorithmic Description:

- Pseudocode vs. Implementation

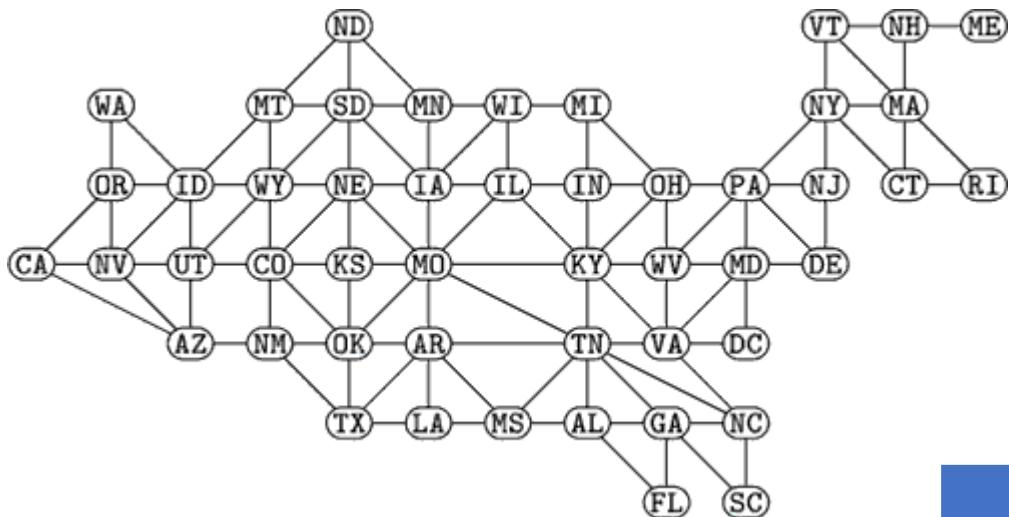
## Algorithm Analysis:

- Time complexity
  - How long does it take to find a solution?
- Space complexity
  - How much memory does it need to perform search?
- Soundness
  - When a solution is returned, is it guaranteed to be correct?
- Completeness
  - Is the algorithm guaranteed to find a solution when one exists?

# Takeaways

- Most decision-making tasks can be formulated as a state space search
- State space search is formalized using graphs, simple paths, search trees, and pseudo code
- DFS and BFS are framed are generic search strategies
- Cycle detection is required to achieve efficiency  
(and completeness in DFS)

# PSet 1



	Visited List	Time (s)	# Paths Popped from Queue	Max Queue Size	Returned Path's Length/Cost
DFS	No				
	Yes				
BFS	No				
	Yes				
IDS	No				
	Yes				